

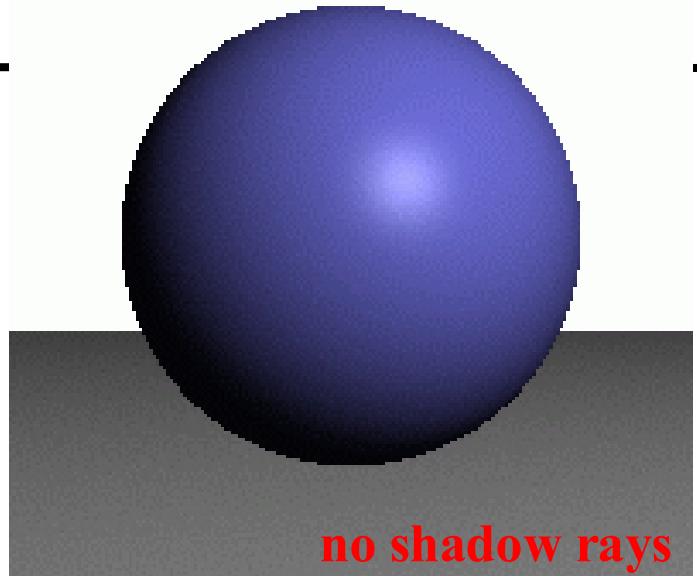
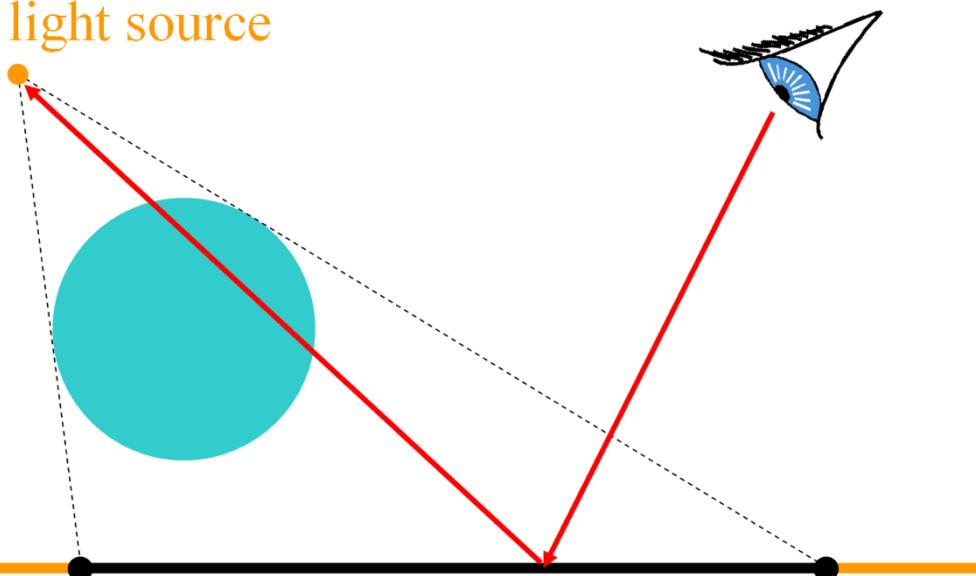
Acceleration Data Structures for Ray Tracing

Most slides are taken from Fredo Durand

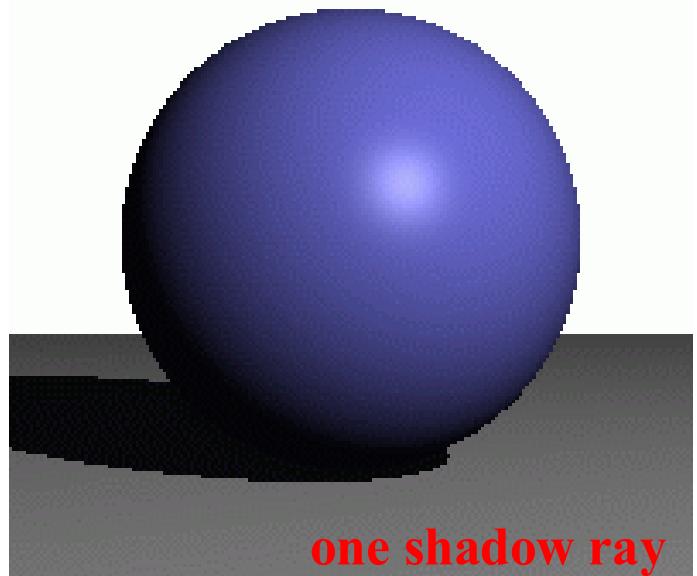
Shadows

- one shadow ray per intersection per point light source

point light source



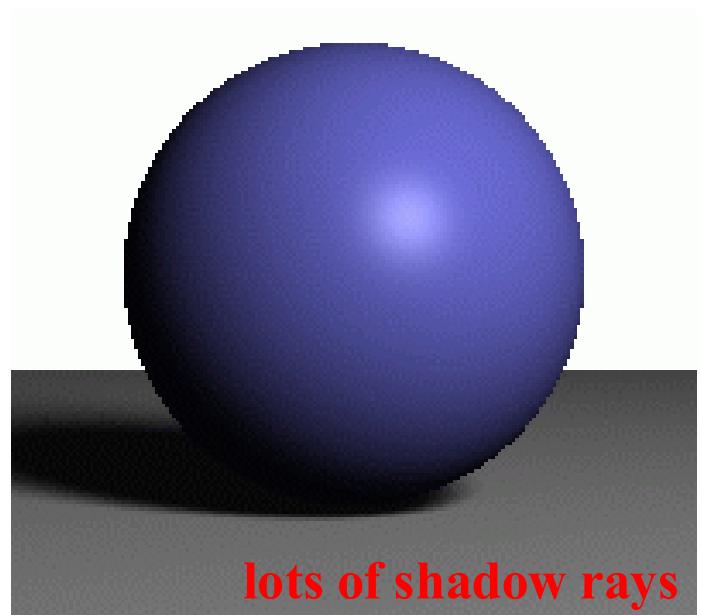
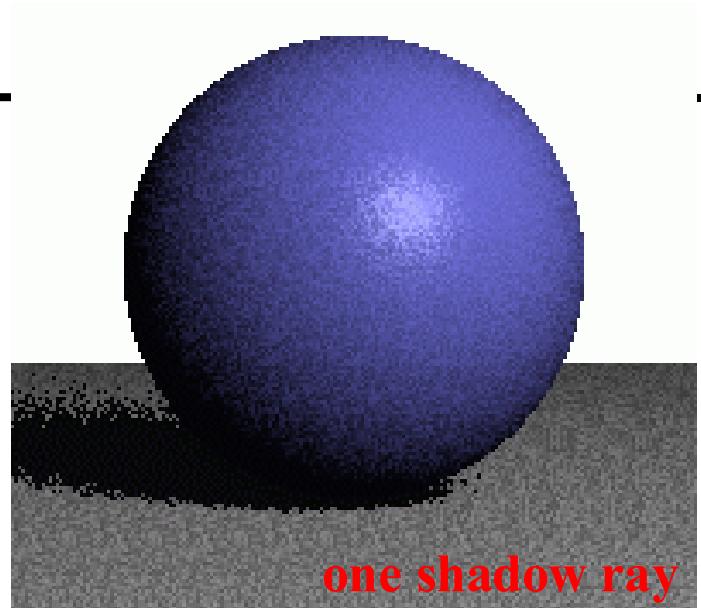
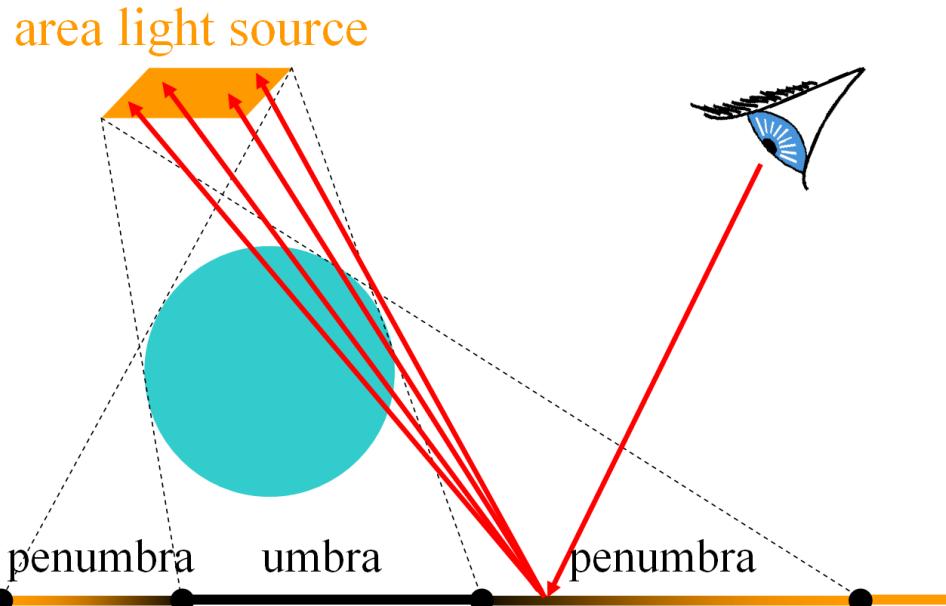
no shadow rays



one shadow ray

Soft Shadows

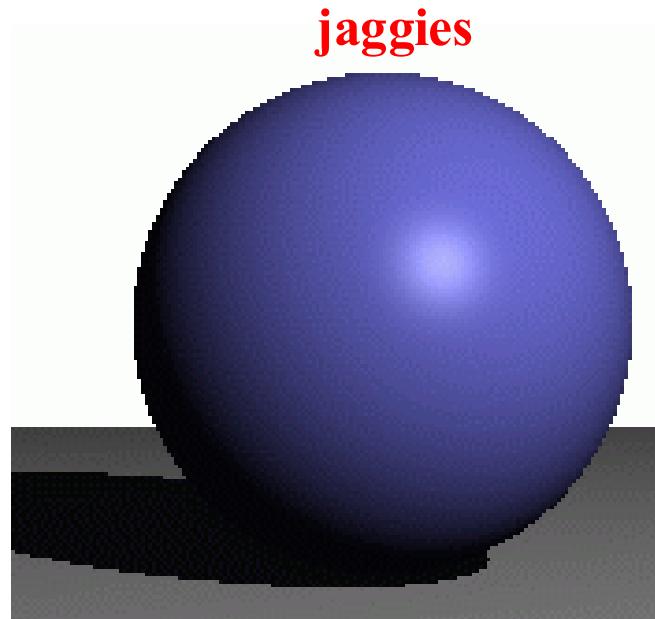
- multiple shadow rays
to sample area light
source



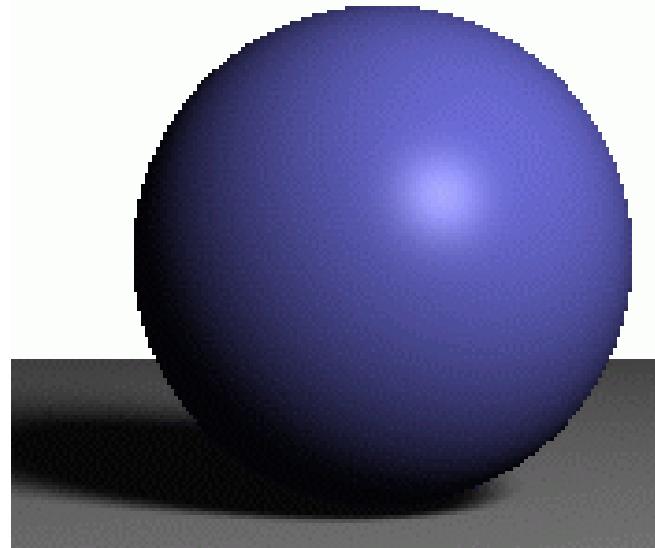
Antialiasing – Supersampling

- multiple rays per pixel

point light

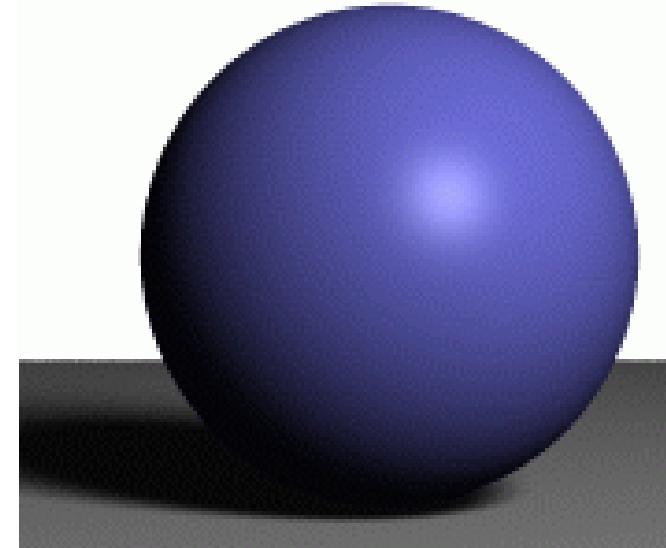
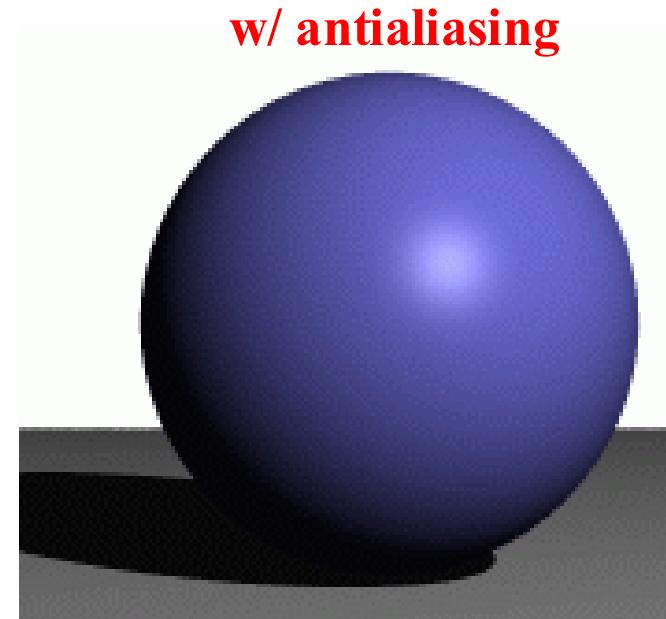


area light



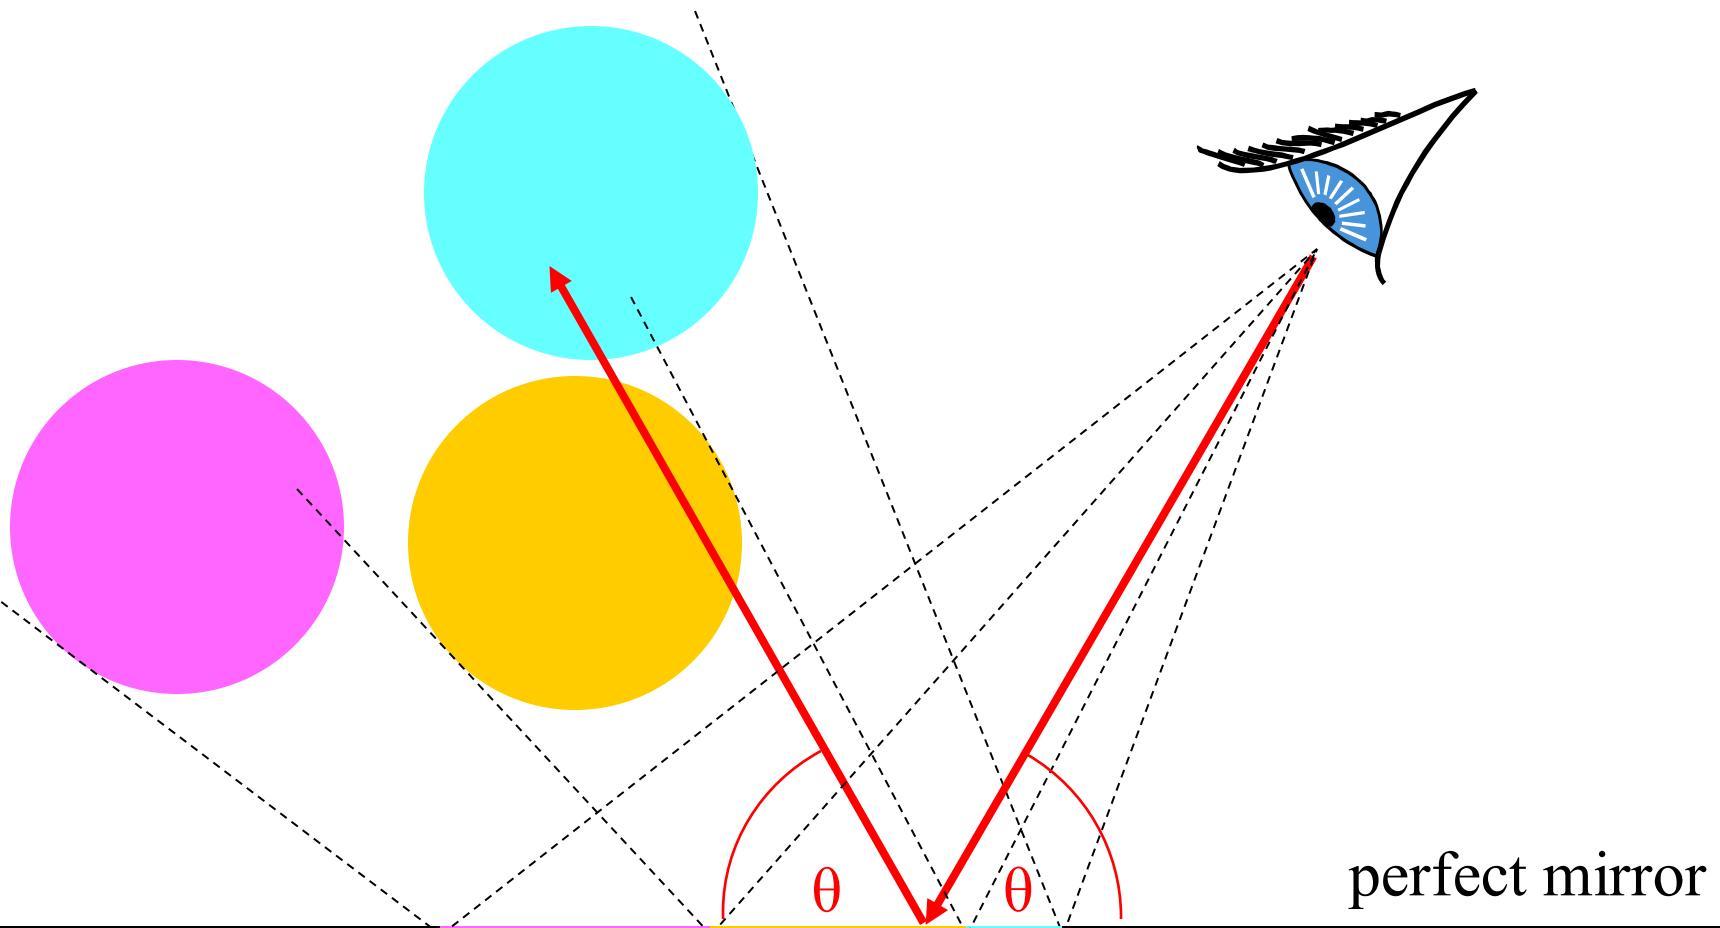
jaggies

w/ antialiasing



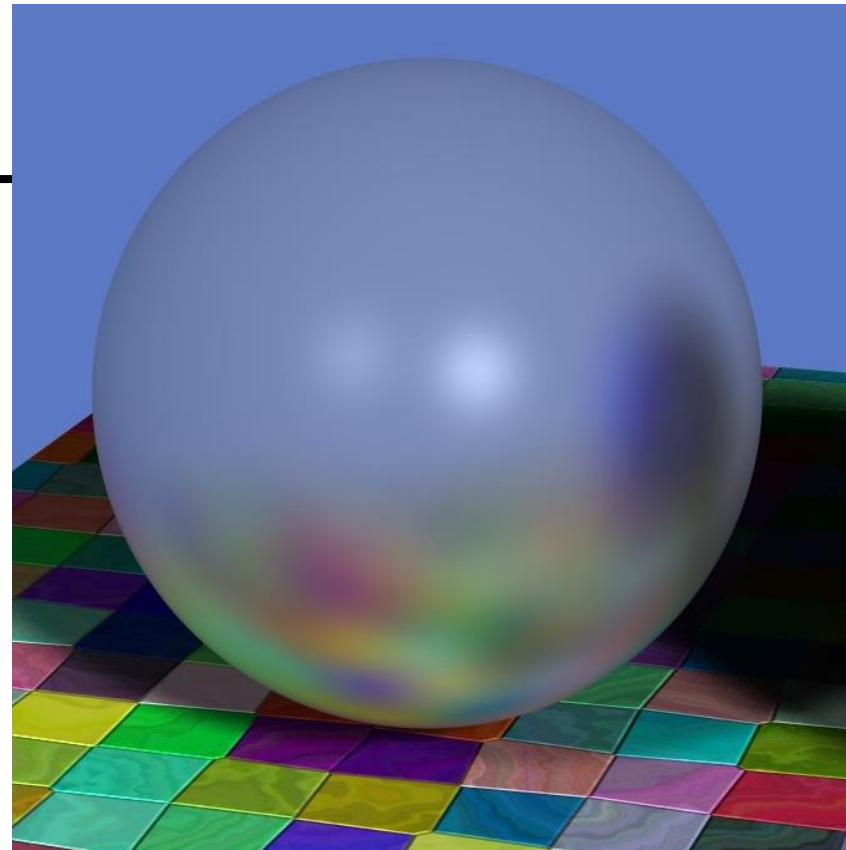
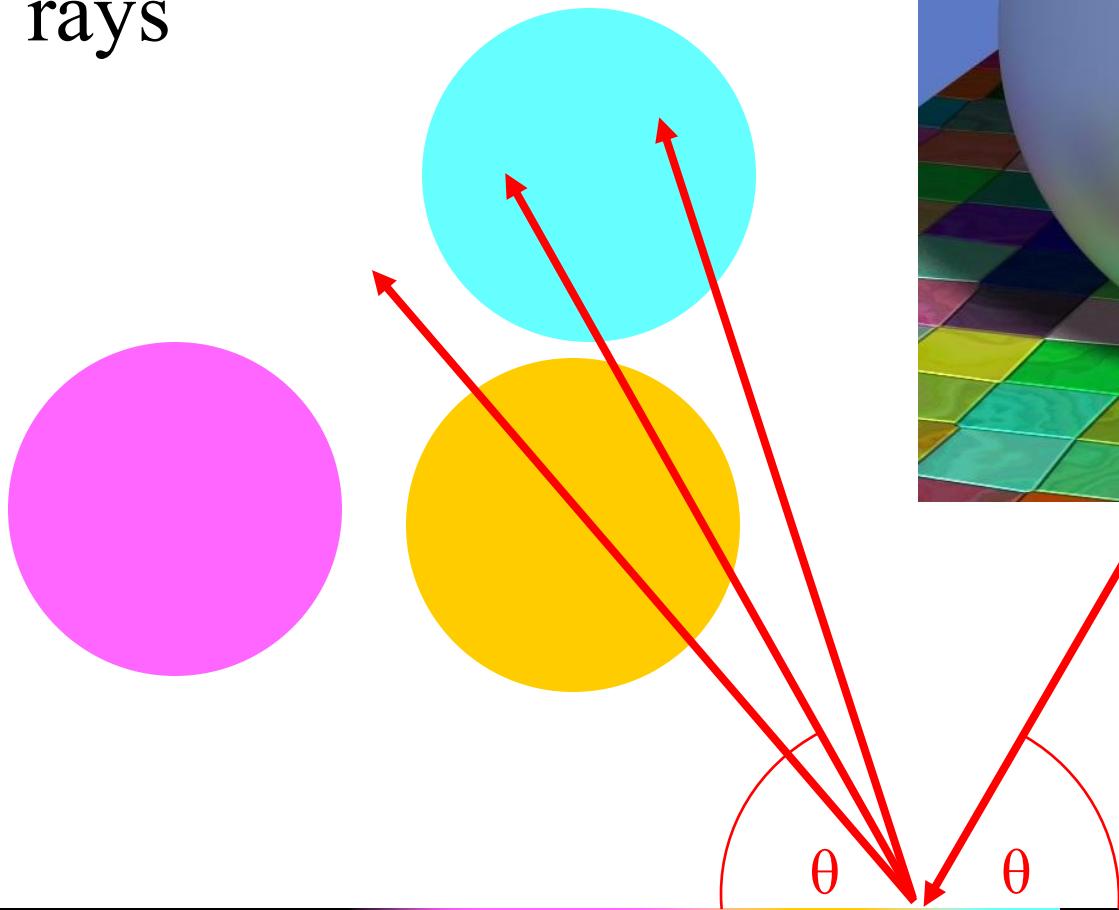
Reflection

- one reflection ray per intersection



Glossy Reflection

- multiple reflection rays

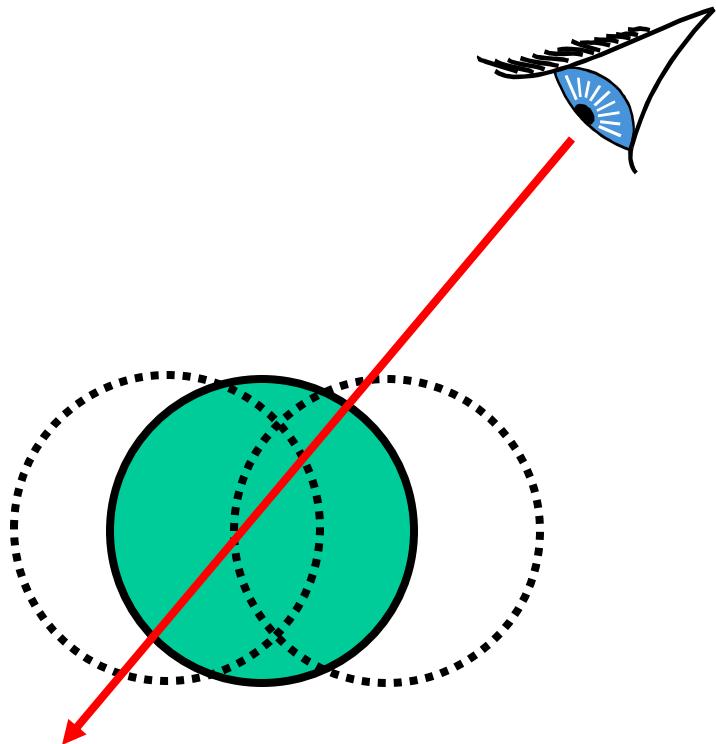


Justin Legakis

polished surface

Motion Blur

- Sample objects temporally



Rob Cook

Algorithm Analysis

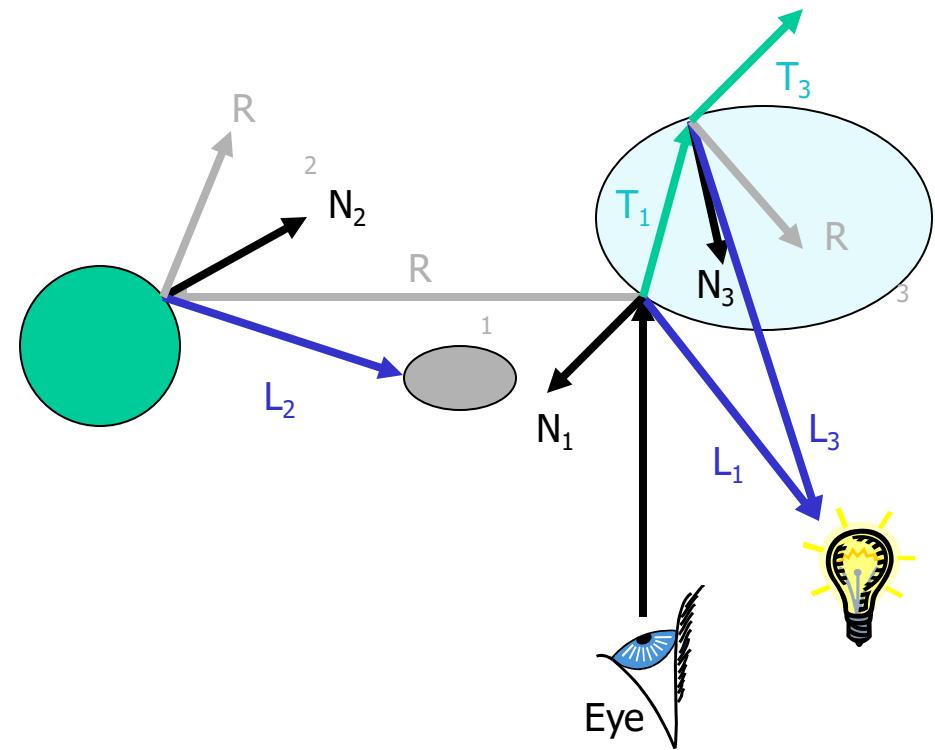
- Ray casting
- Lots of primitives
- Recursive
- Distributed Ray Tracing Effects
 - Soft shadows
 - Anti-aliasing
 - Glossy reflection
 - Motion blur
 - Depth of field

$\text{cost} \leq \text{height} * \text{width} * \text{num primitives} *$
intersection cost *
num shadow rays *
super-sampling *
num glossy rays *
num temporal samples *
 $3^{\max \text{ recursion depth}} *$
...

can we reduce this?

The cost of Ray Tracing

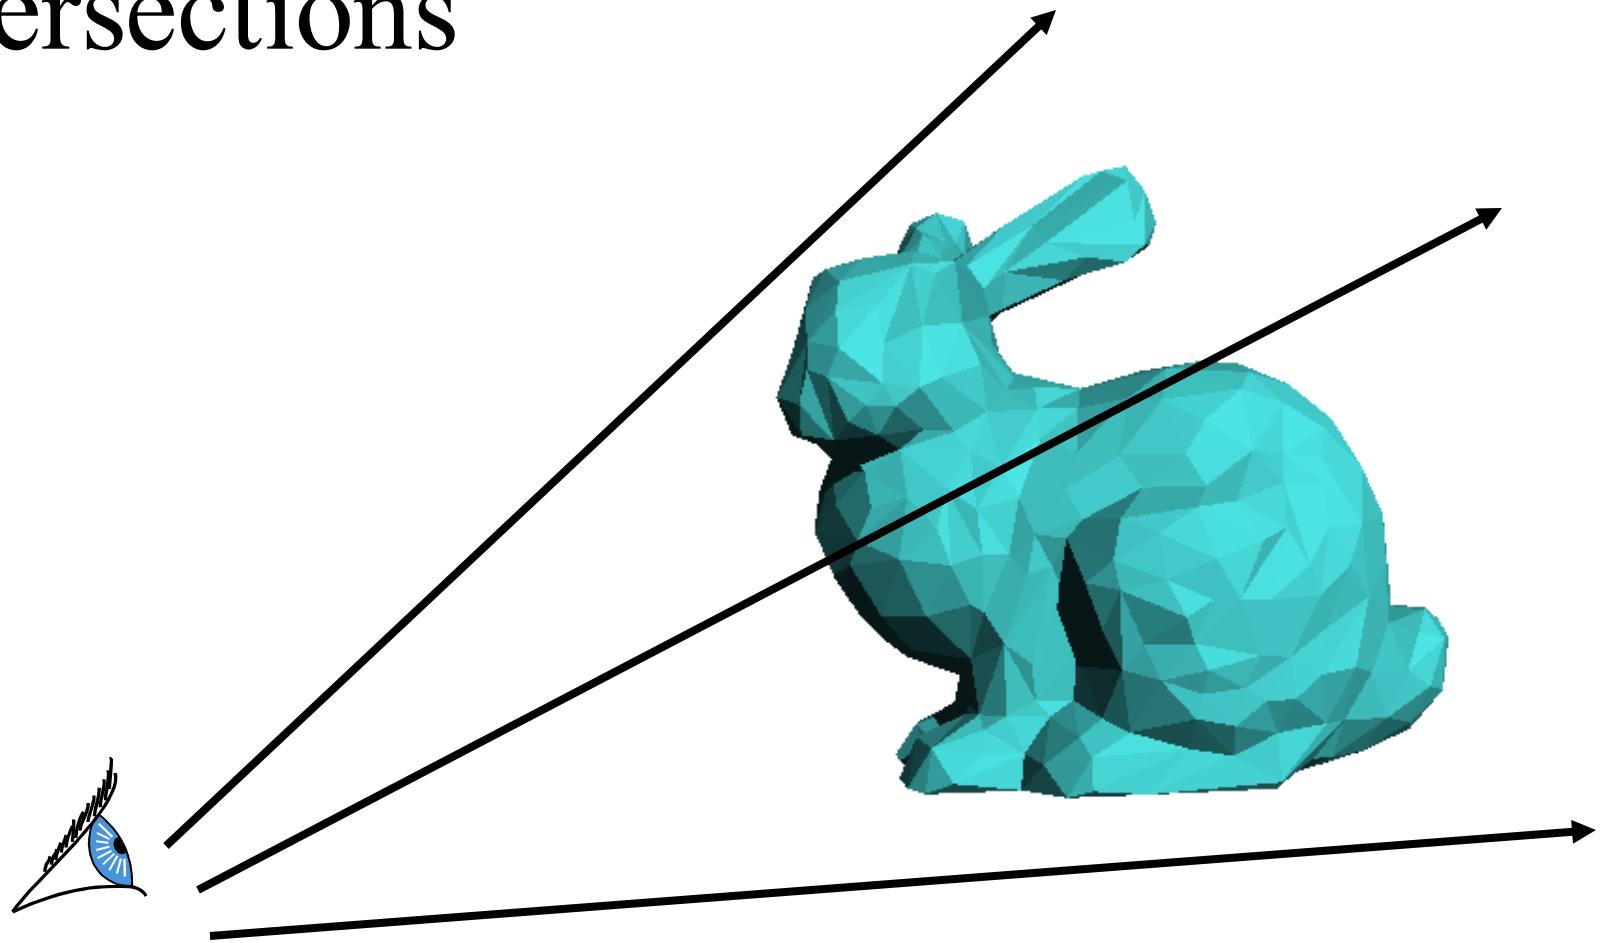
- Many Primitives
 - Many Rays
 - Expensive Intersections



Accelerating Ray Tracing

- Four main groups of acceleration techniques:
 - Reducing the average cost of intersecting a ray with a scene:
 - Faster intersection calculations
 - Fewer intersection calculations
 - Reducing the total number of rays that are traced
 - Adaptive recursion depth control
 - Parallelization, specialized hardware

Reduce the number of ray/primitive intersections

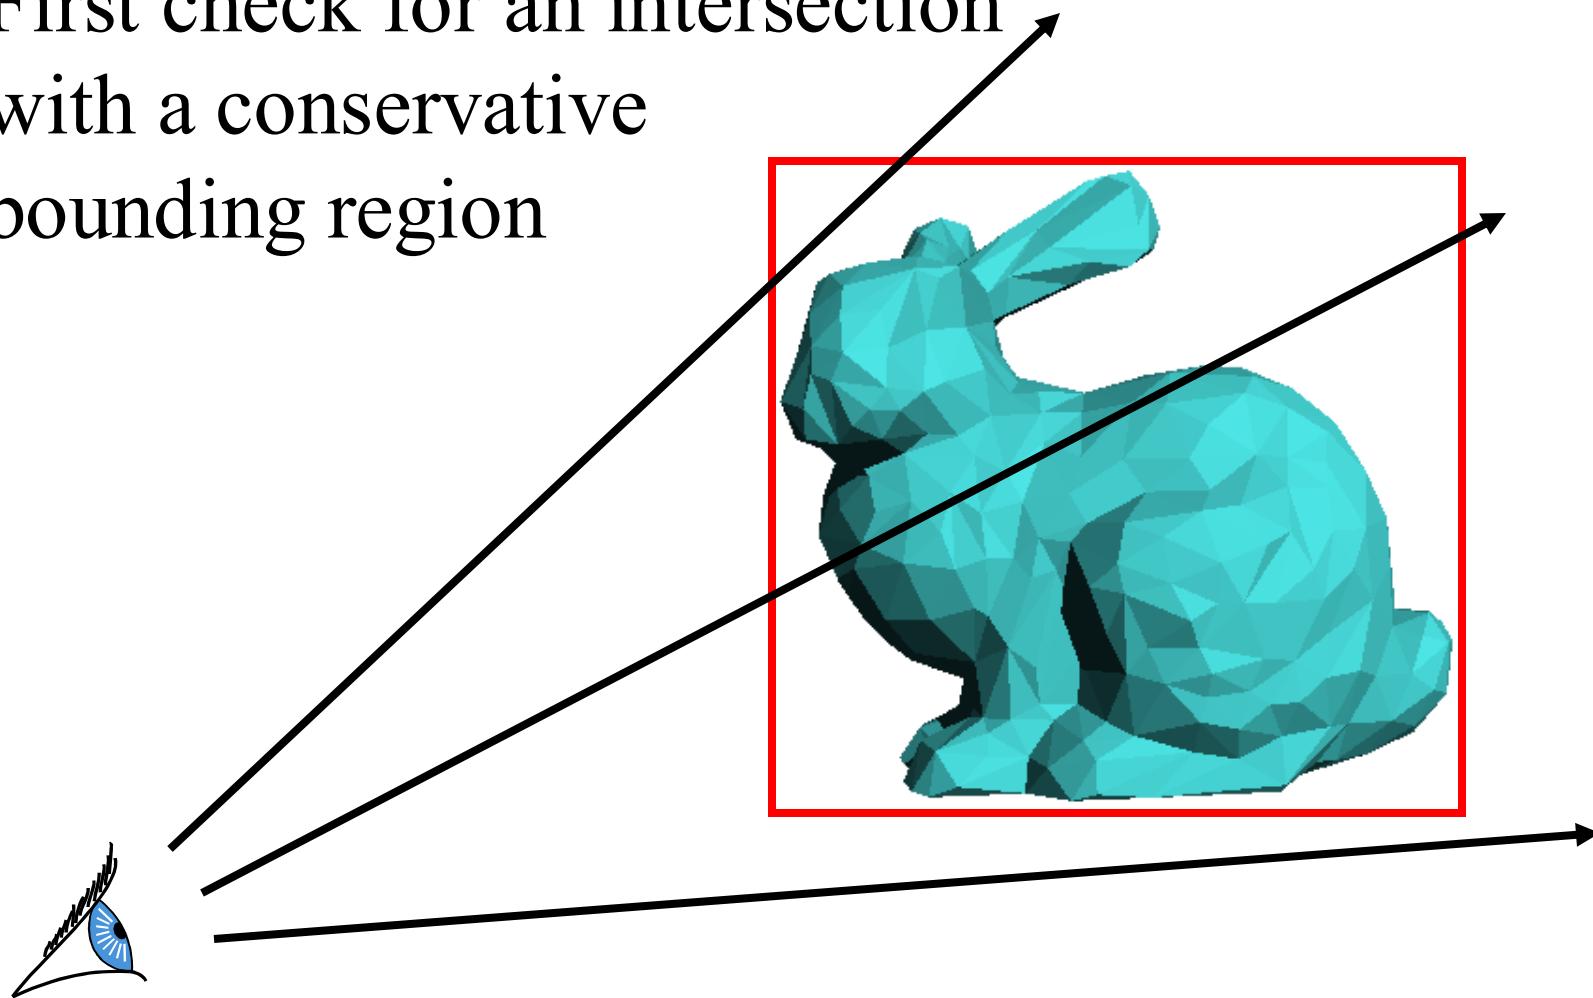


Bounding Volumes

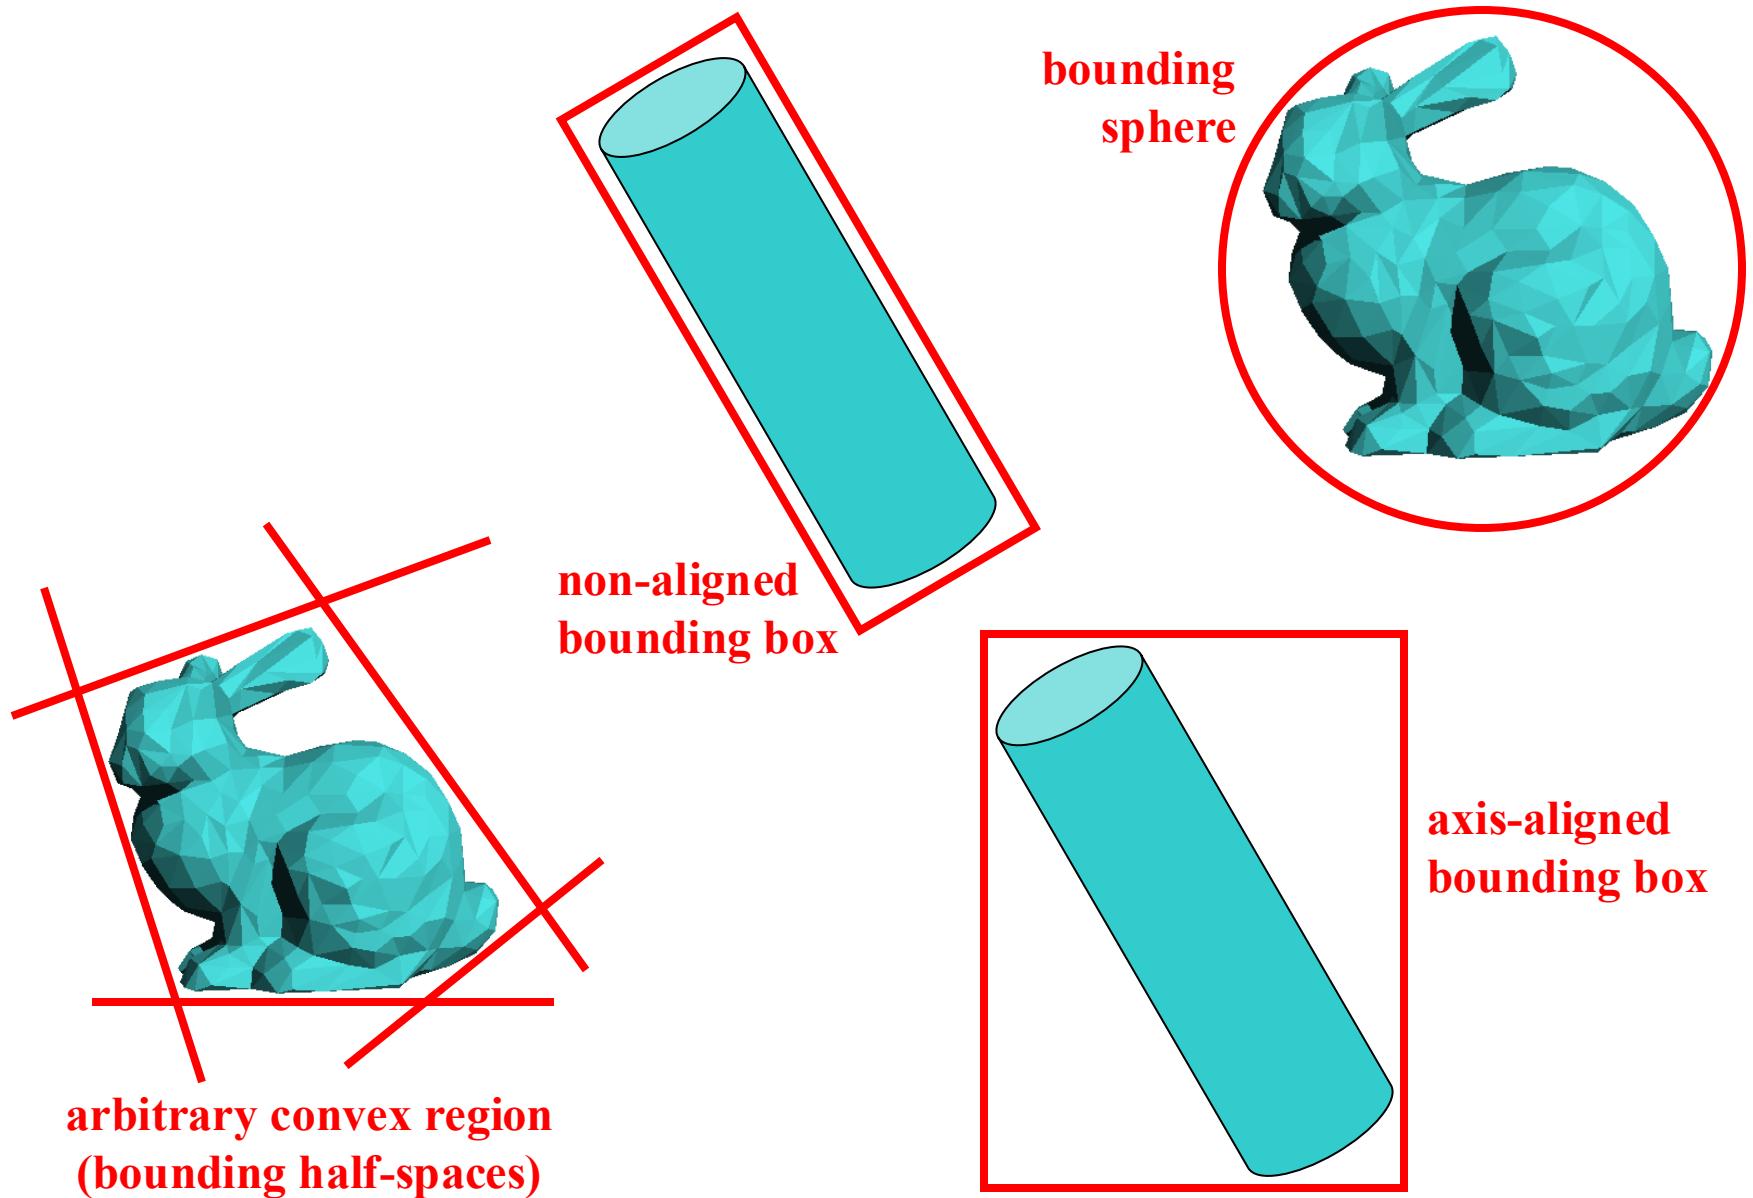
- Idea: associate with each object a simple bounding volume. If a ray misses the bounding volume, it also misses the object contained therein.
- Effective for additional applications:
 - Clipping acceleration
 - Collision detection

Early reject

- First check for an intersection with a conservative bounding region

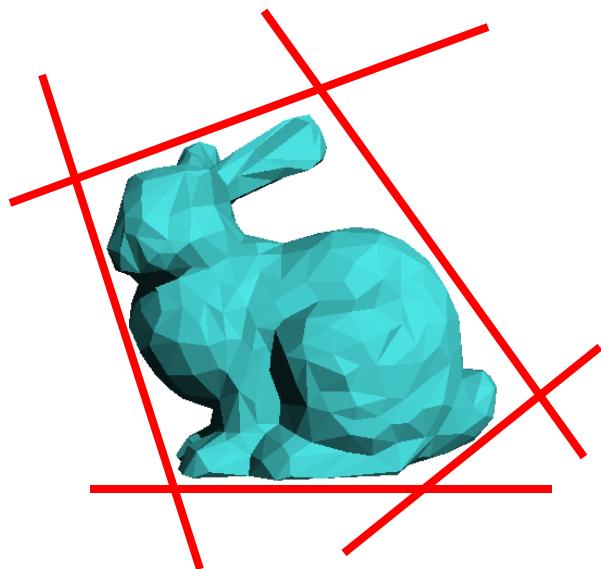


Conservative Bounding Regions

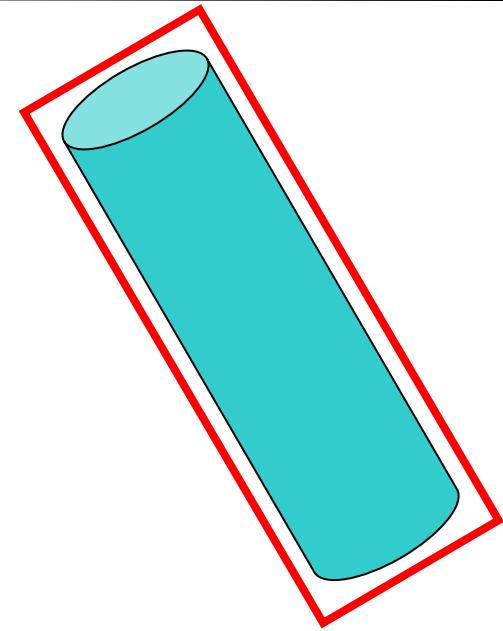
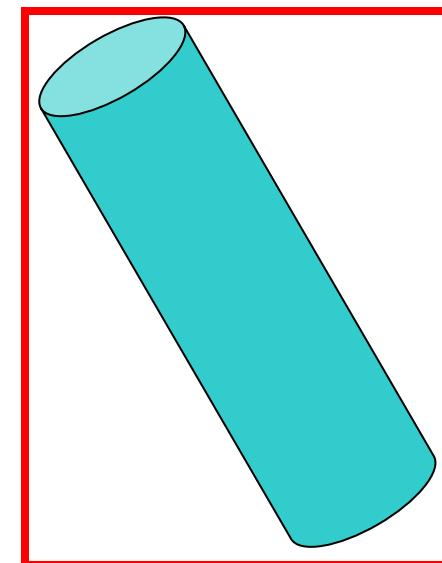


What is a good bounding volume?

- tight → avoid false positives
- fast to intersect
- easy to construct



arbitrary convex region
(bounding half-spaces)



Bounding Volumes

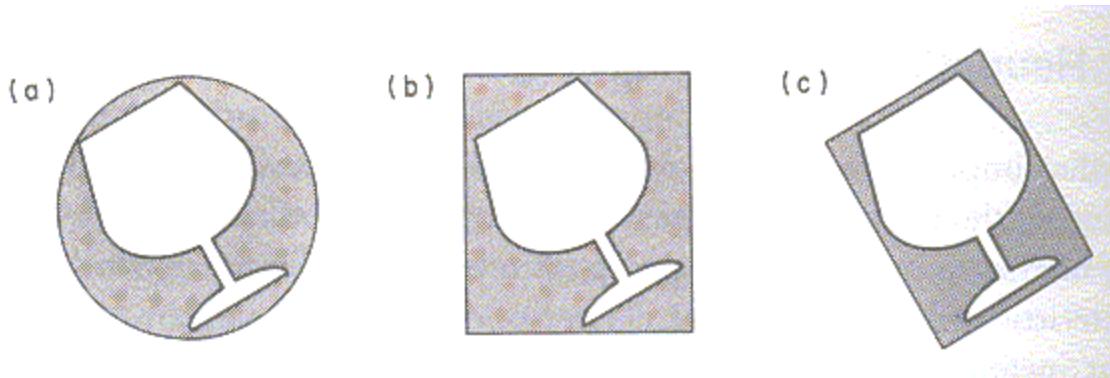
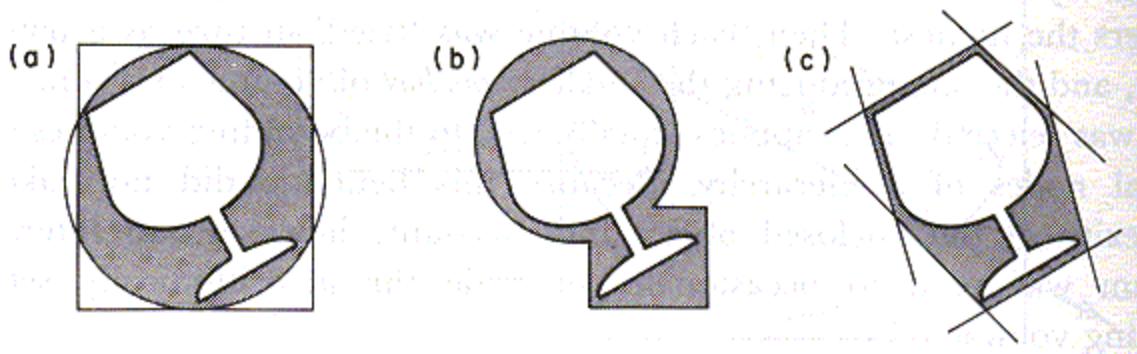
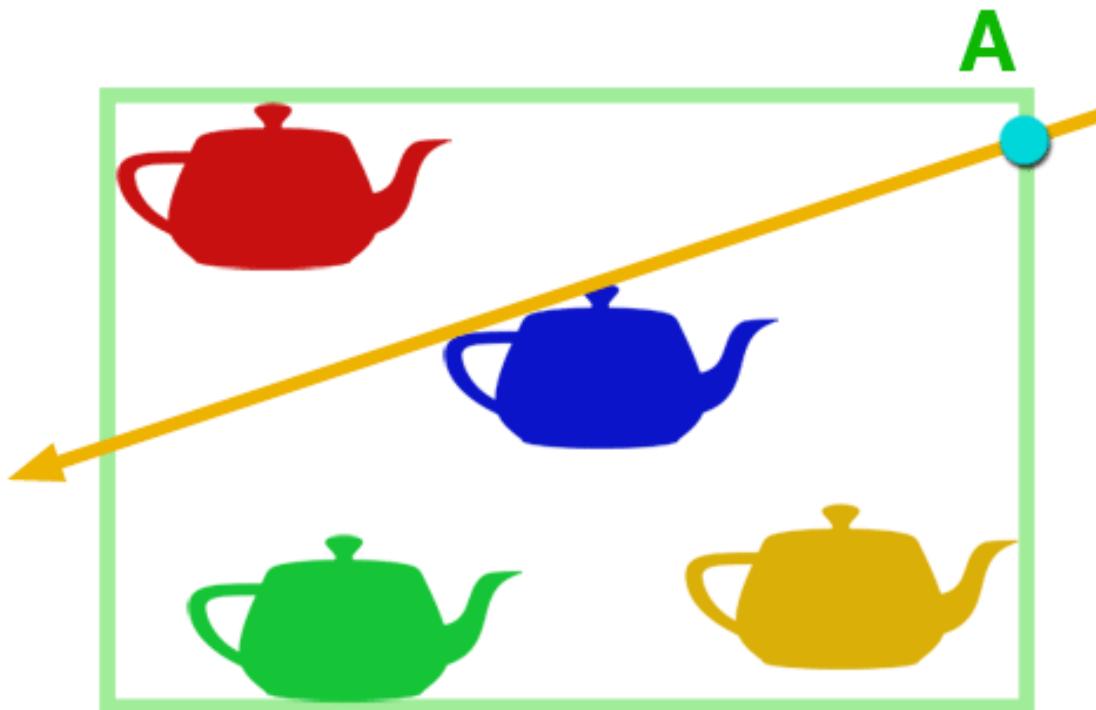


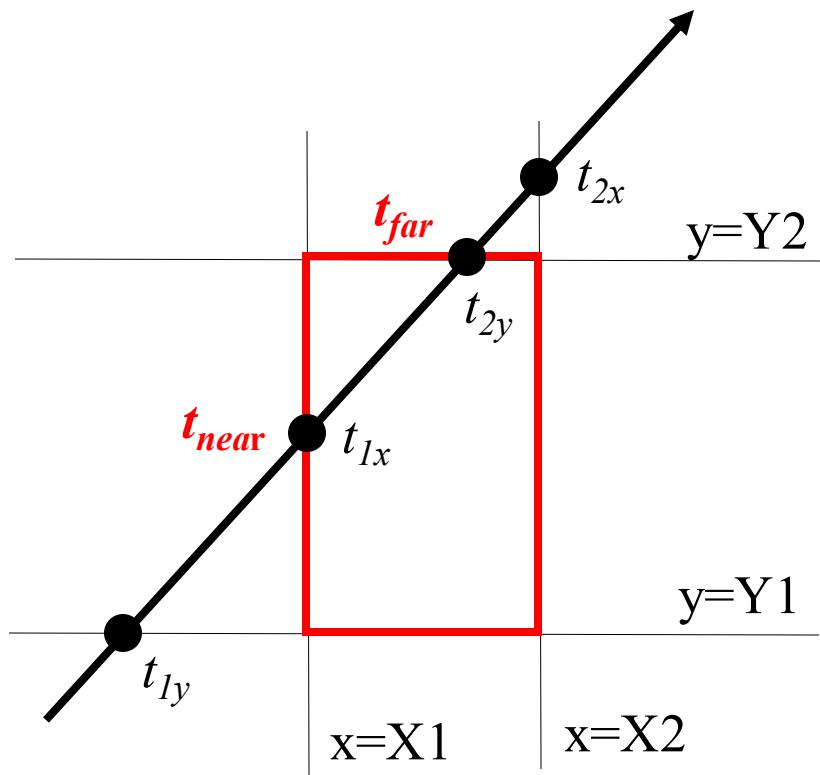
Figure 1.10: Three different bounding volumes for the same teacup model.



Hierarchical Bounding Boxes



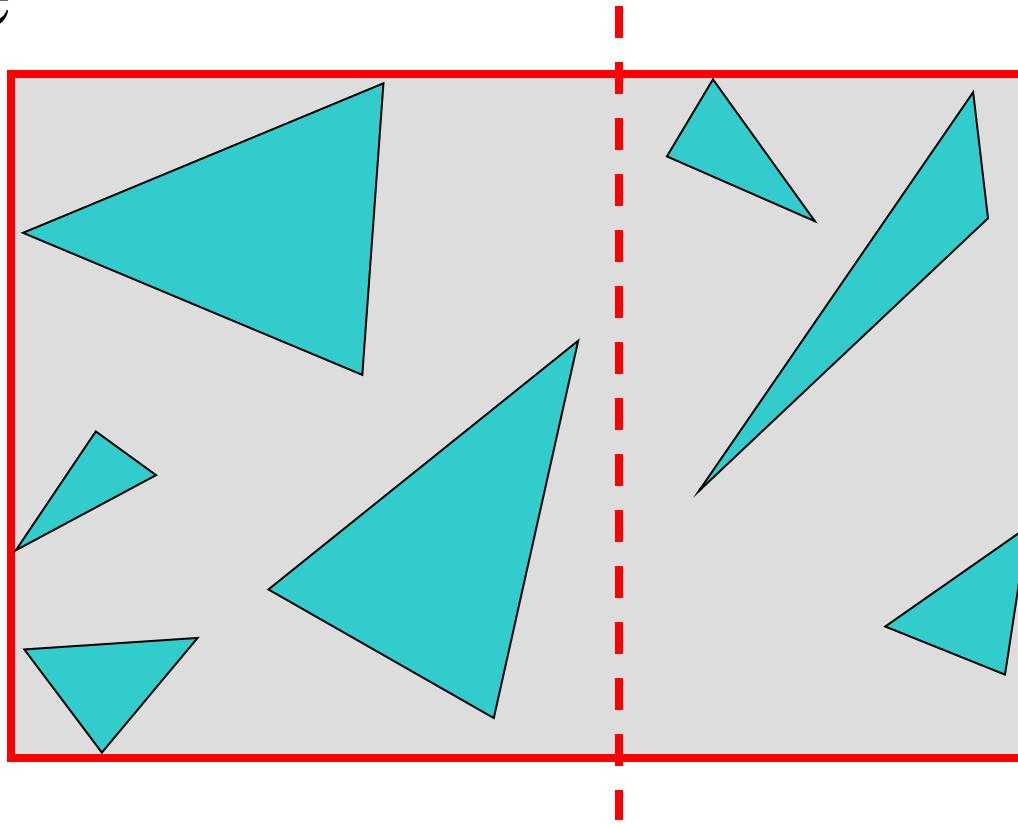
Intersection with Axis-Aligned Box



- For all 3 axes, calculate the intersection distances t_l and t_2
- $t_{near} = \max(t_{lx}, t_{ly}, t_{lz})$
 $t_{far} = \min(t_{2x}, t_{2y}, t_{2z})$
- If $t_{near} > t_{far}$, box is missed
- If $t_{far} < t_{min}$, box is behind
- If box survived tests, report intersection at t_{near}

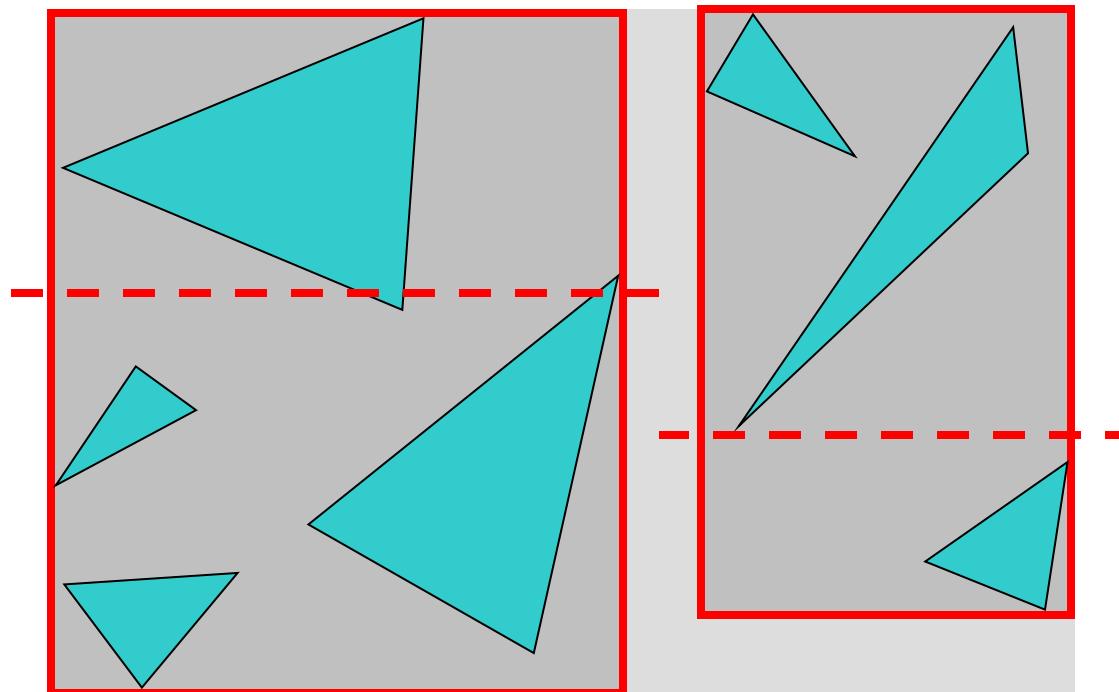
Bounding Volume Hierarchy

- Find bounding box of objects
- Split objects into two groups
- Recurse



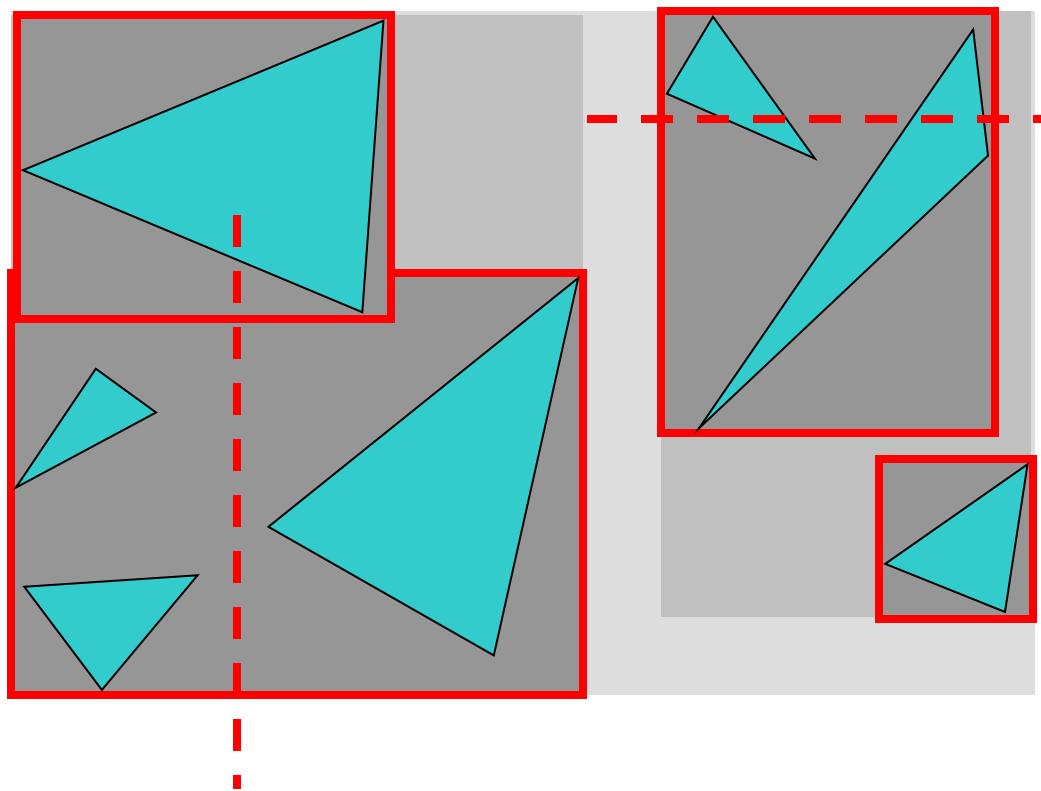
Bounding Volume Hierarchy

- Find bounding box of objects
- Split objects into two groups
- Recurse

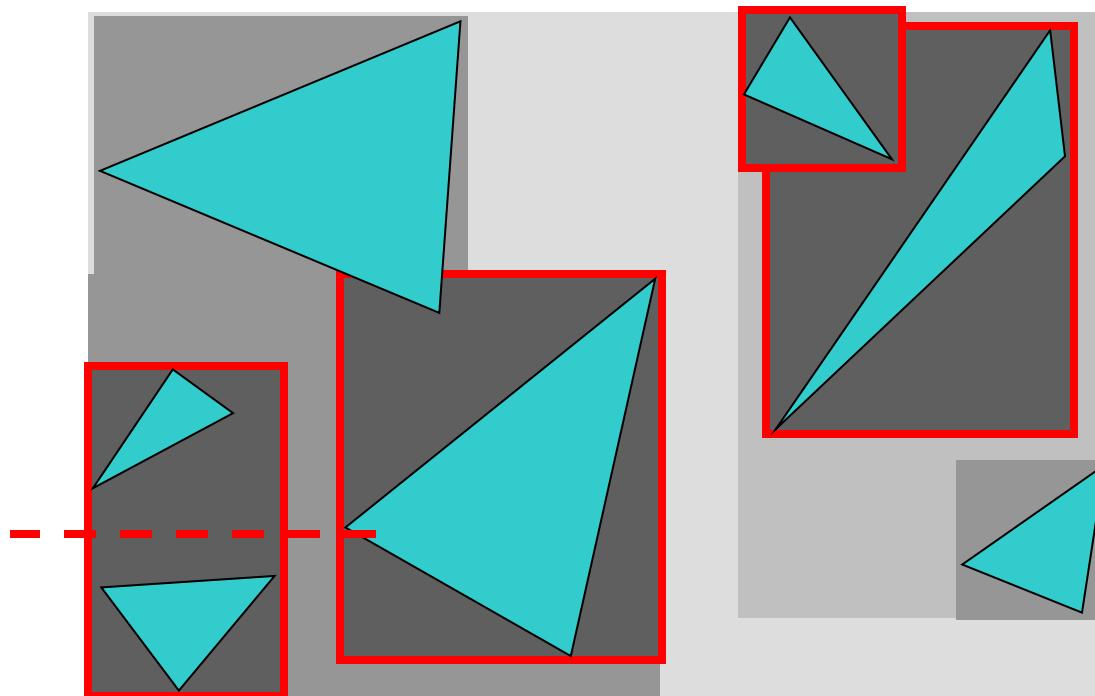


Bounding volumes can intersect

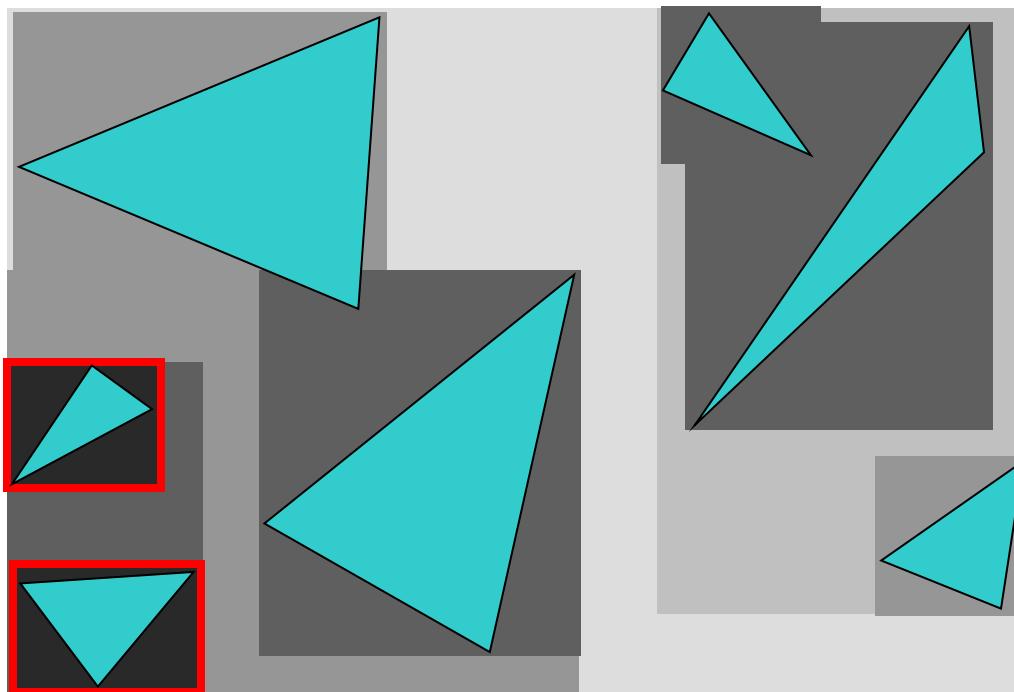
- Find bounding box of objects
- Split objects into two groups
- Recurse



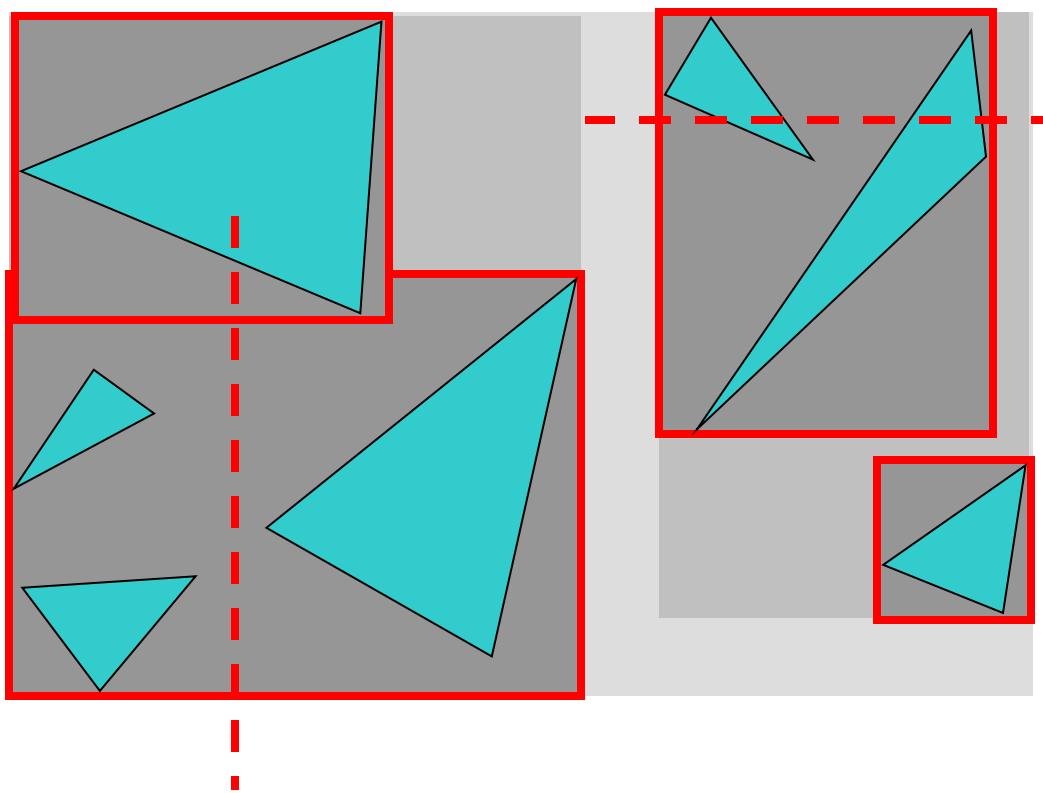
Bounding Volume Hierarchy



Bounding Volume Hierarchy

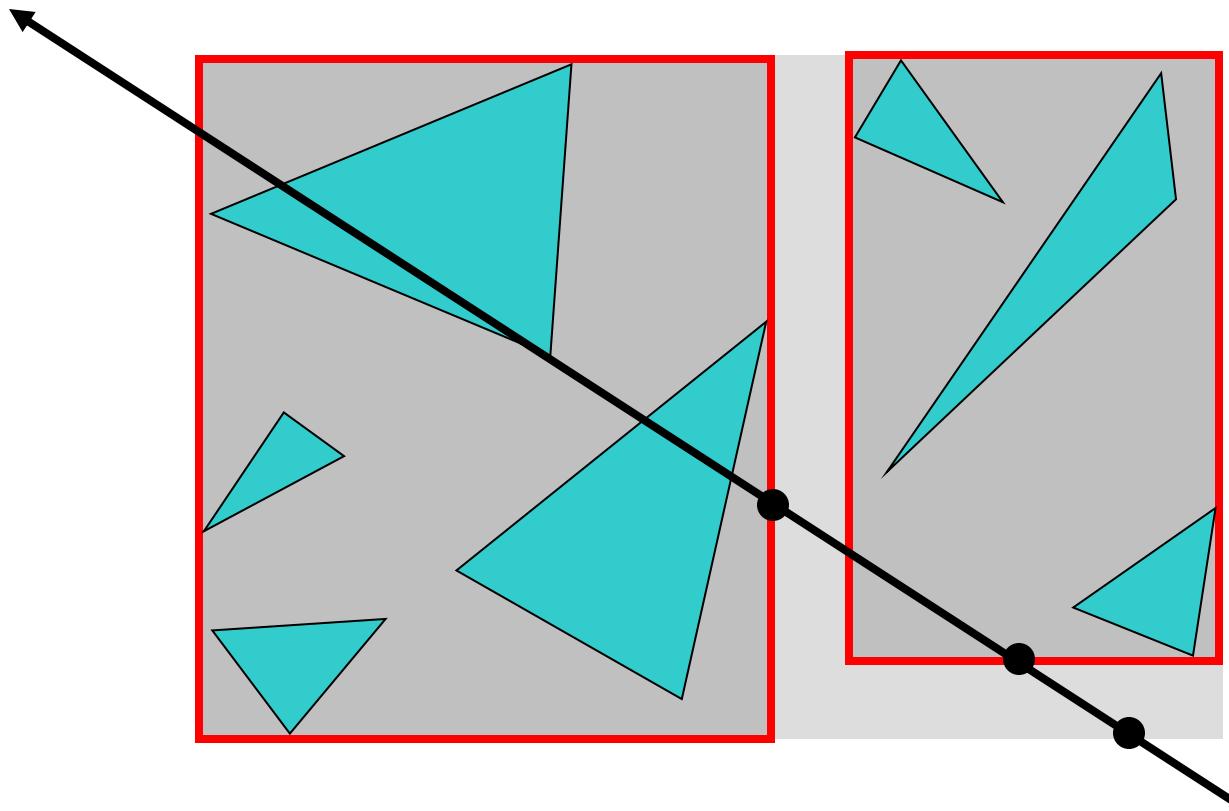


Where to split objects?



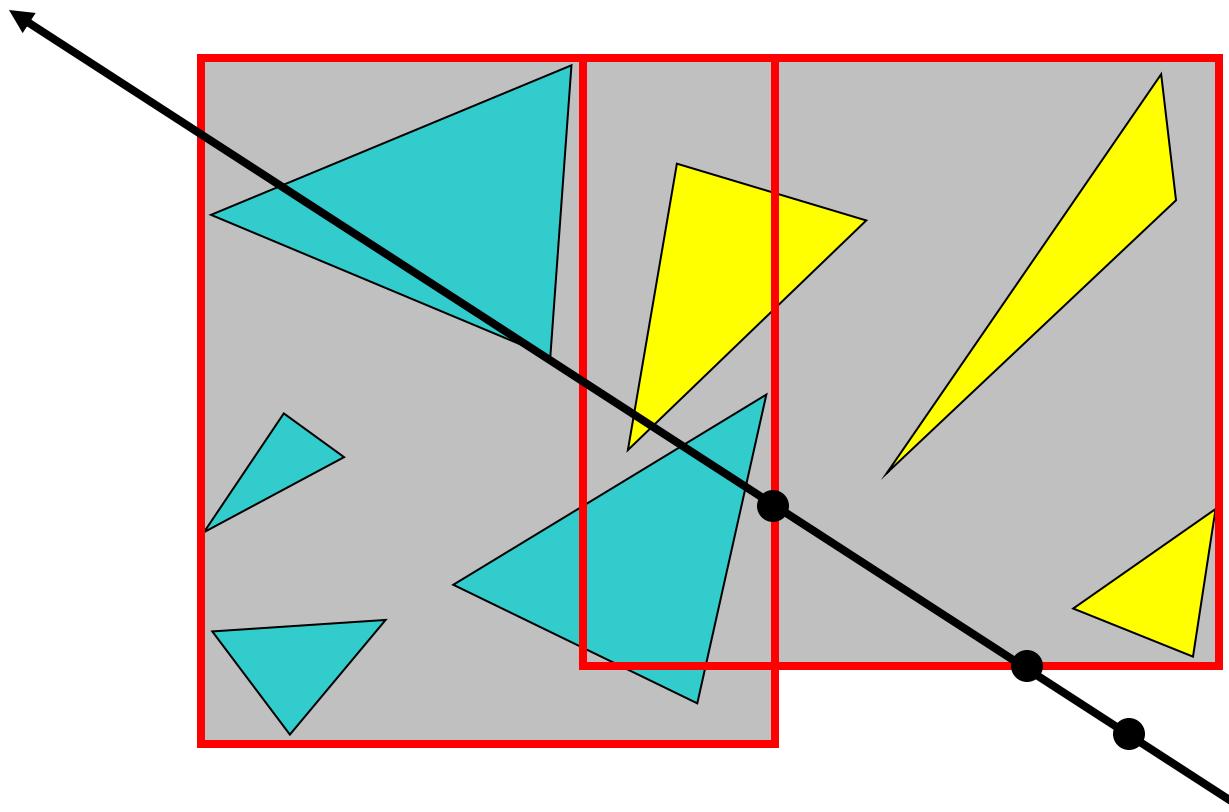
Intersection with BVH

- Check subvolume with closer intersection first

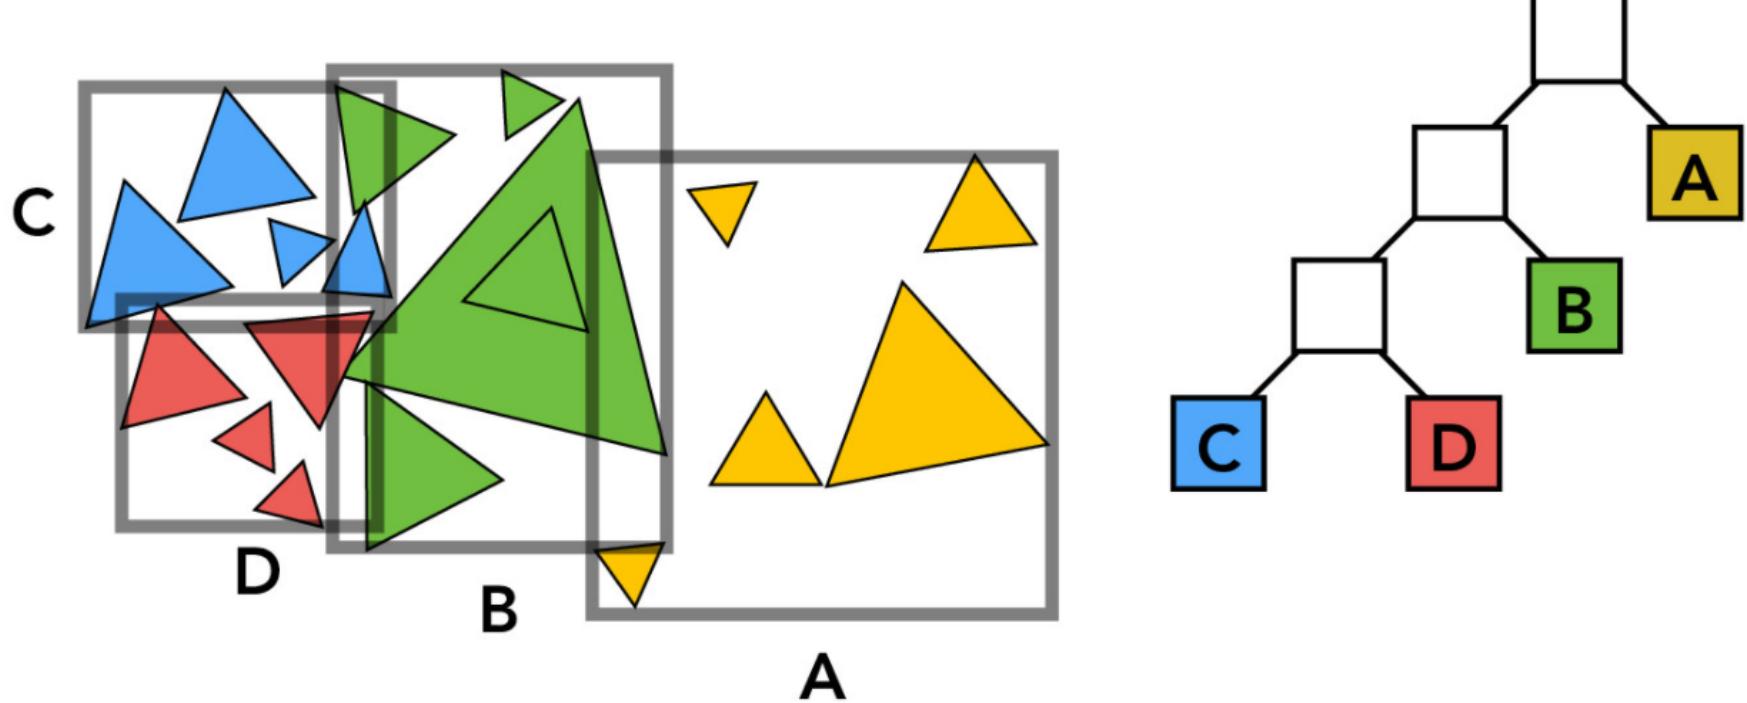


Intersection with BVH

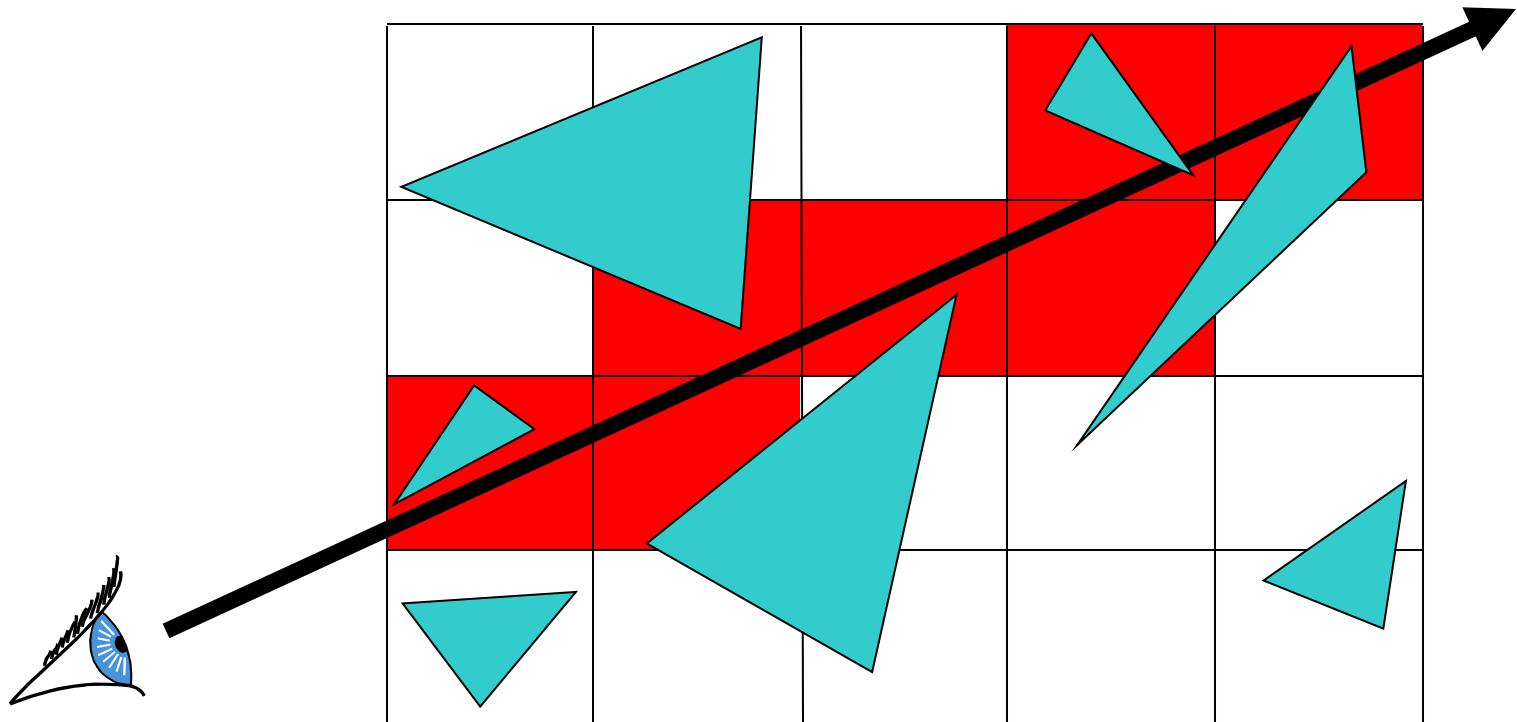
- Don't return intersection immediately if the other subvolume may have a closer intersection



Bounding Volume Hierarchy (BVH)



Spatial Subdivision

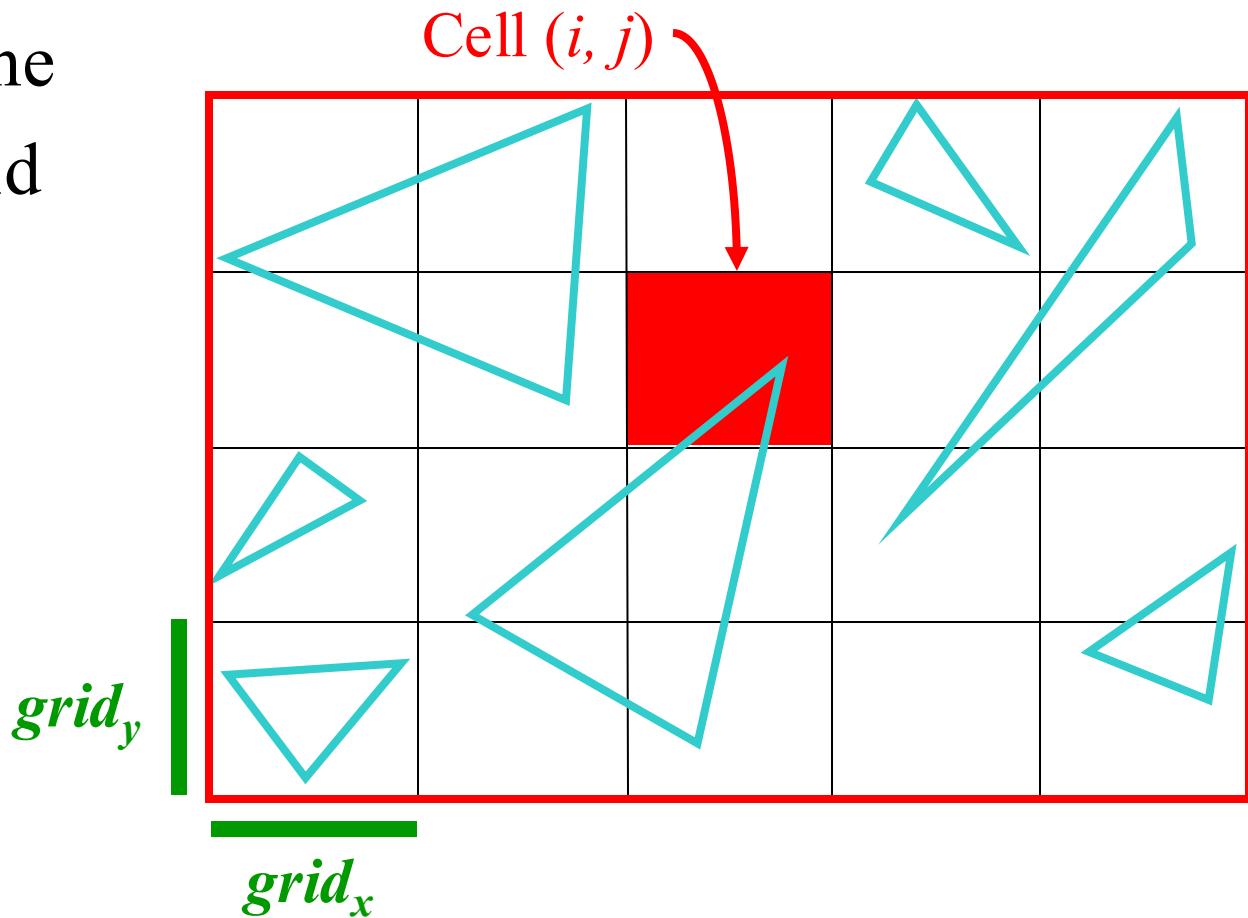


Spatial Subdivision

- Uniform spatial subdivision:
 - The space containing the scene is subdivided into a uniform grid of cubes “voxels”.
 - Each voxel stores a list of all objects at least partially contained in it.
 - Given a ray, voxels are traversed using a 3D variant of the 2D line drawing algorithms.
 - At each voxel the ray is tested for intersection with the primitives stored therein
 - Once an intersection has been found, there is no need to continue to other voxels.

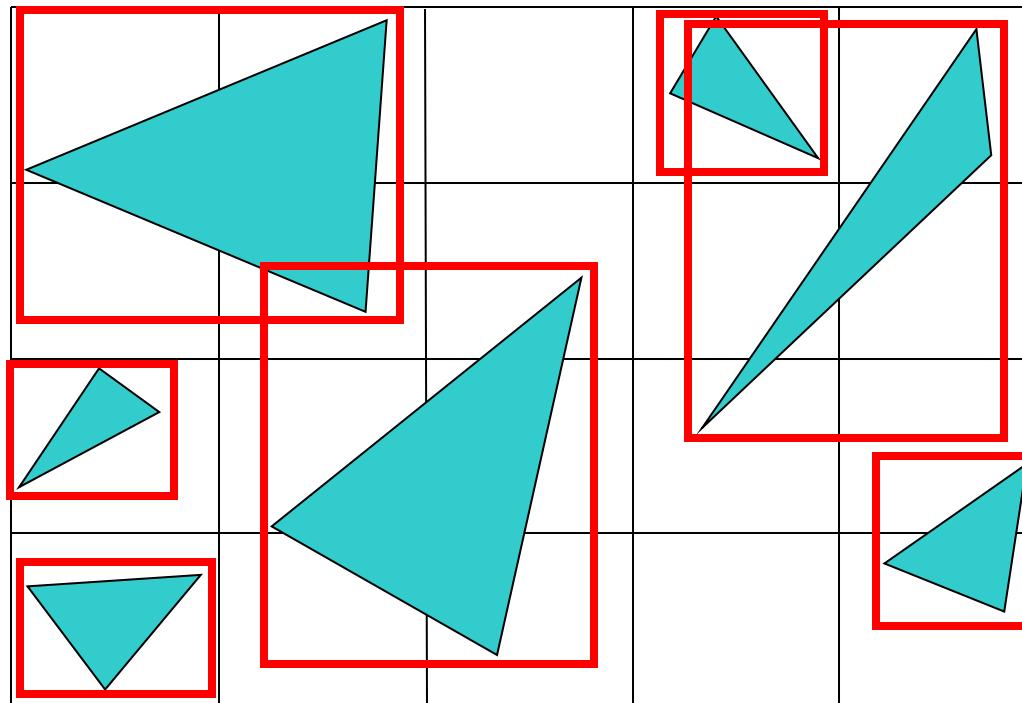
Create grid

- Find bounding box of scene
- Choose grid spacing
- grid_x need not = grid_y



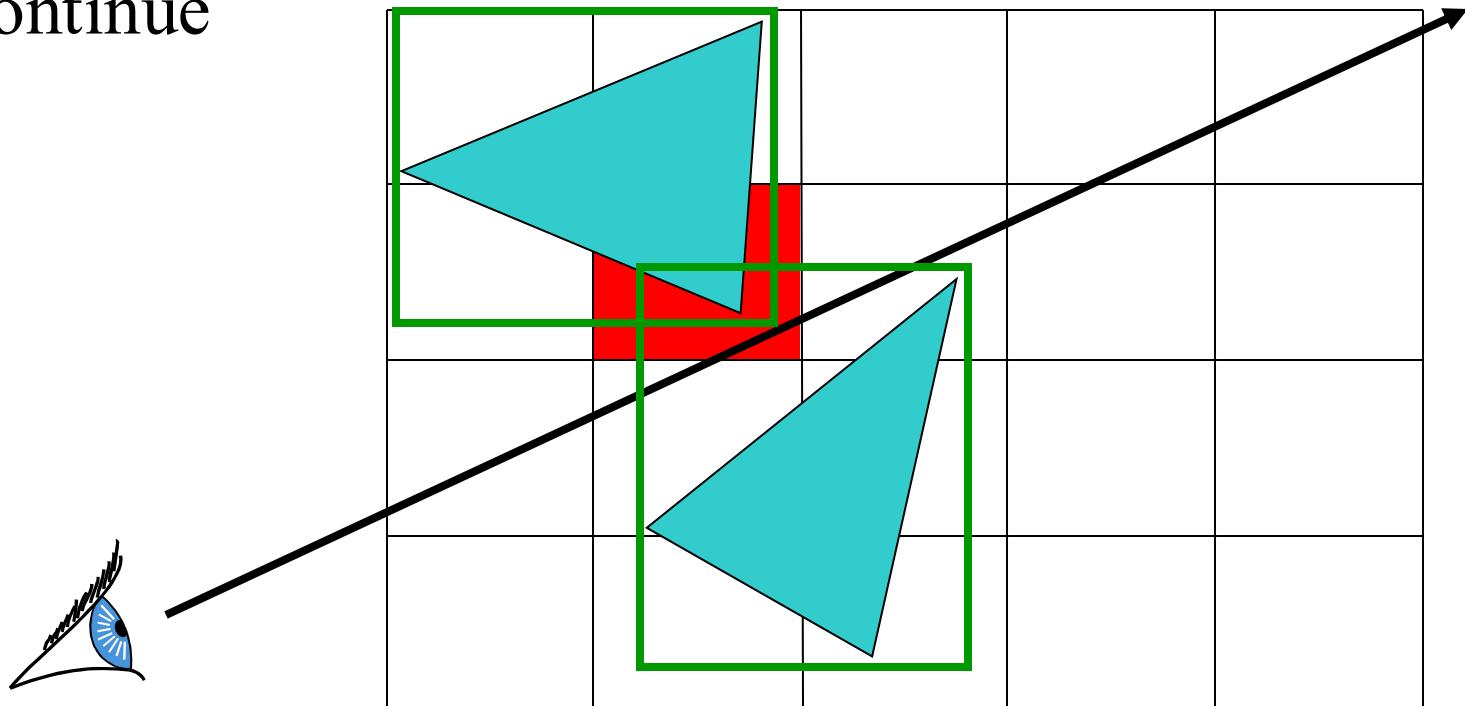
Insert primitives into grid

- Primitives that overlap multiple cells?
- Insert into multiple cells (use pointers)



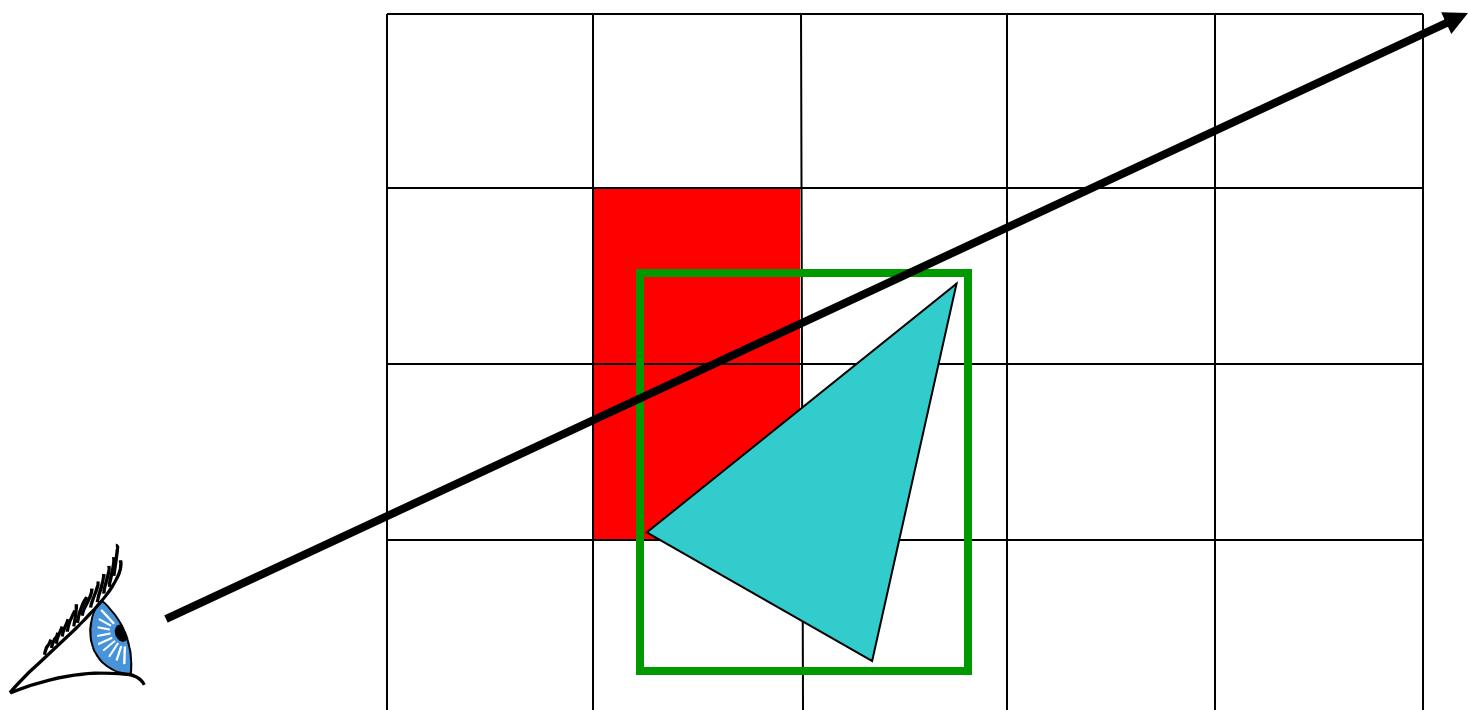
For each cell along a ray

- Does the cell contain an intersection?
- Yes: return closest intersection
- No: continue



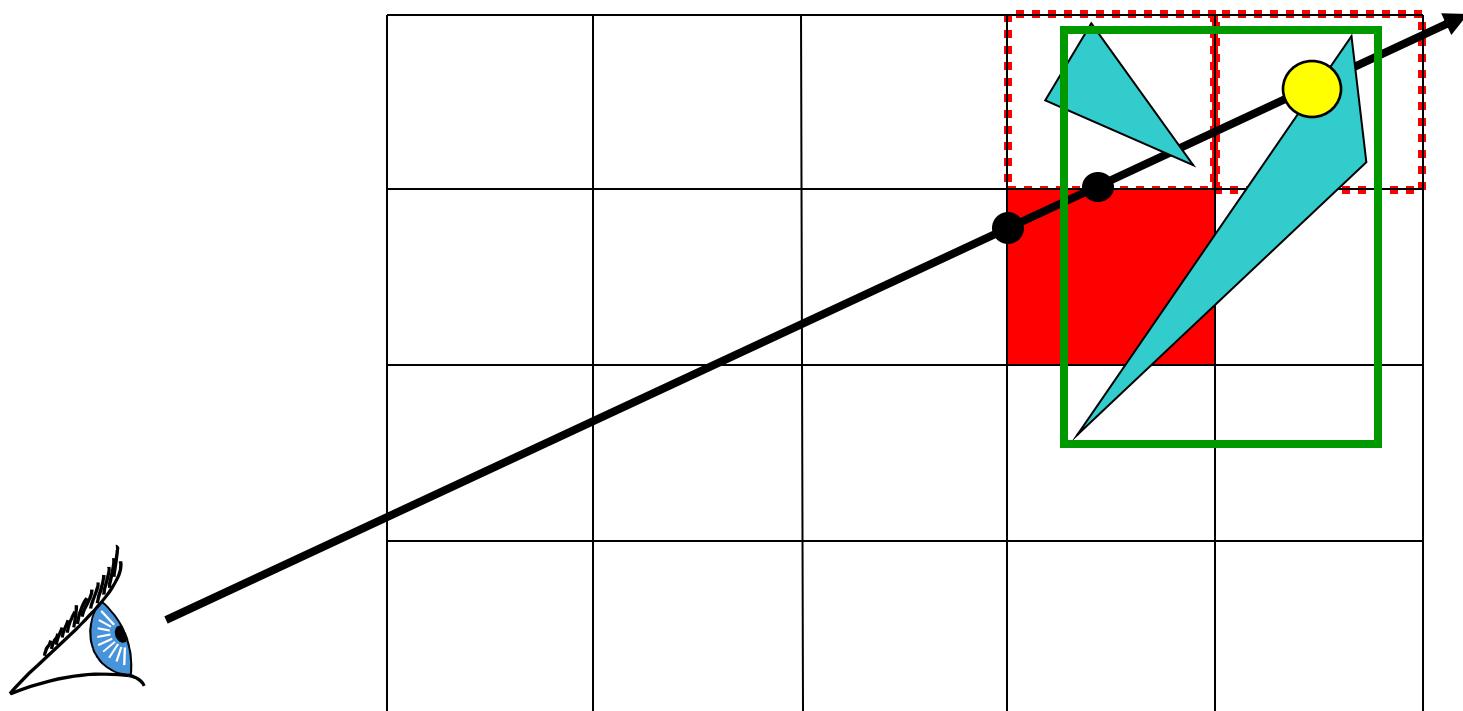
Preventing repeated computation

- Perform the computation once, "mark" the object
- Don't re-intersect marked objects



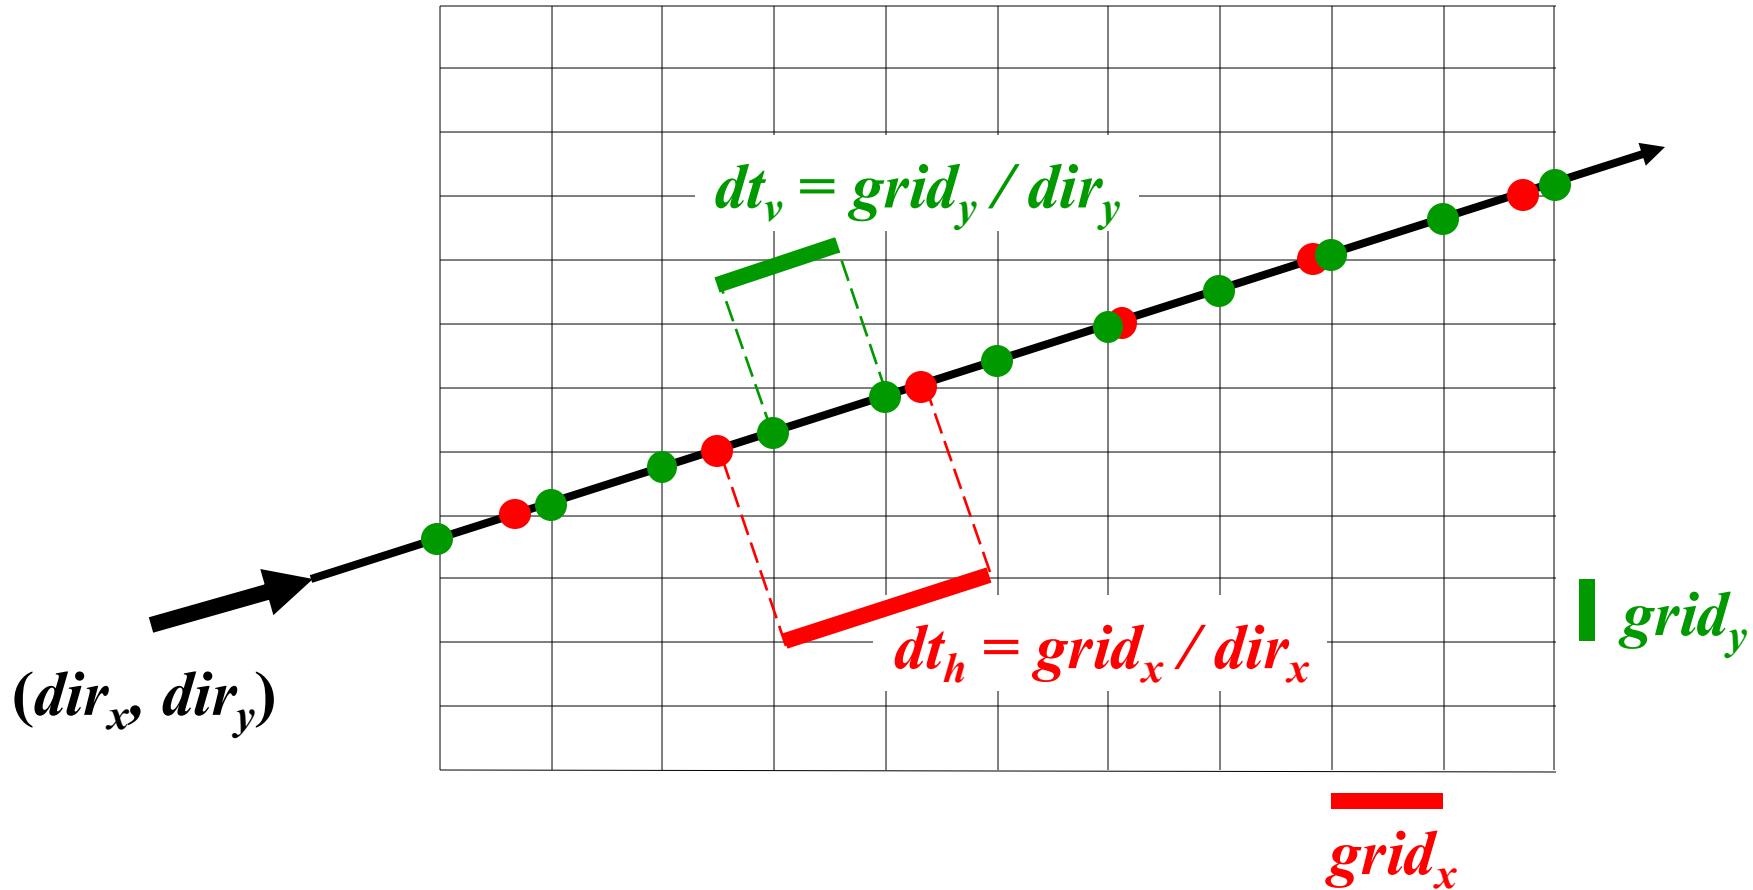
Don't return distant intersections

- If intersection t is not within the cell range, continue
(there may be something closer)



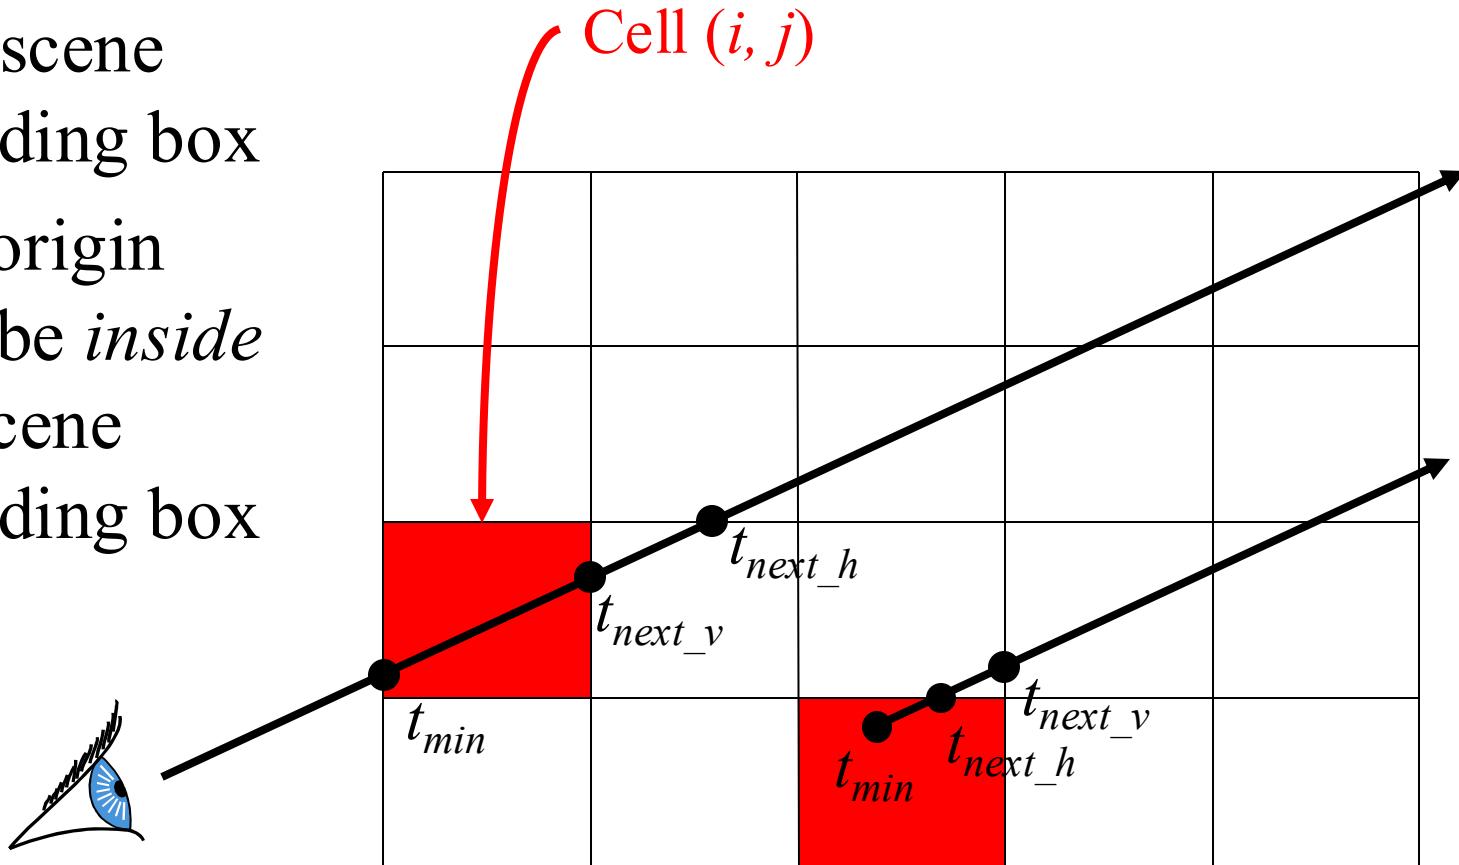
Is there a pattern to cell crossings?

- Yes, the horizontal and vertical crossings have regular spacing



Where do we start?

- Intersect ray with scene bounding box
- Ray origin may be *inside* the scene bounding box



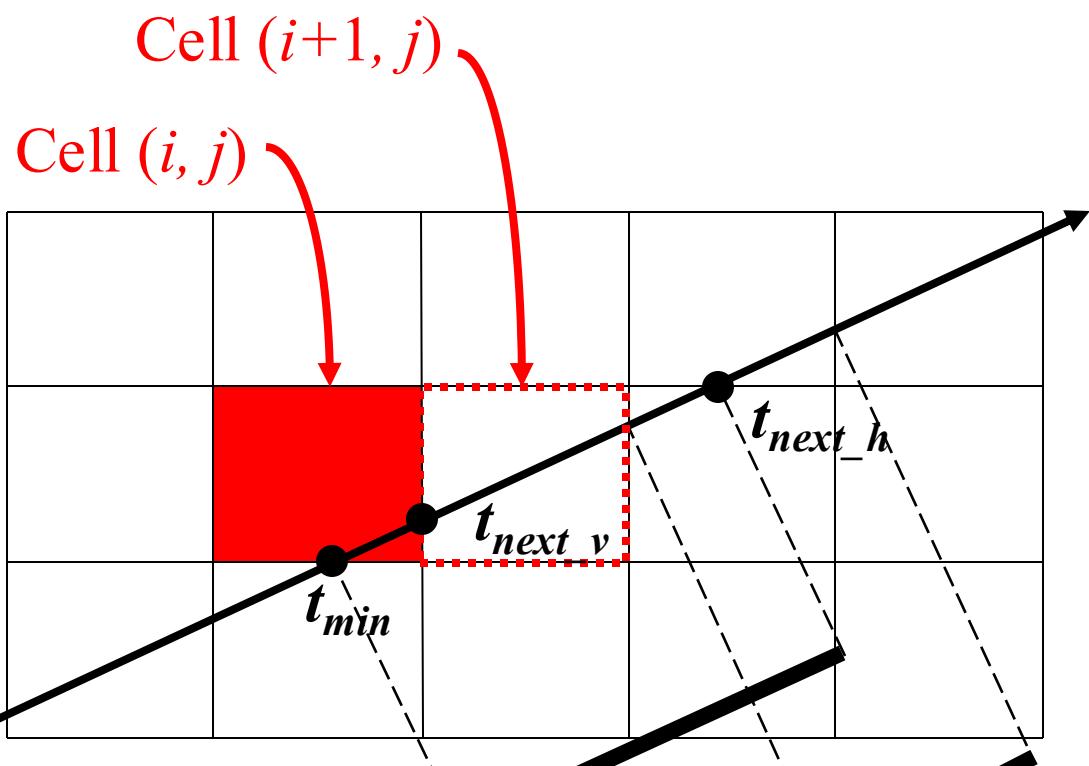
What's the next cell?

```
if  $t_{next\_v} < t_{next\_h}$ 
   $i += sign_x$ 
   $t_{min} = t_{next\_v}$ 
   $t_{next\_v} += dt_v$ 
```

```
else
```

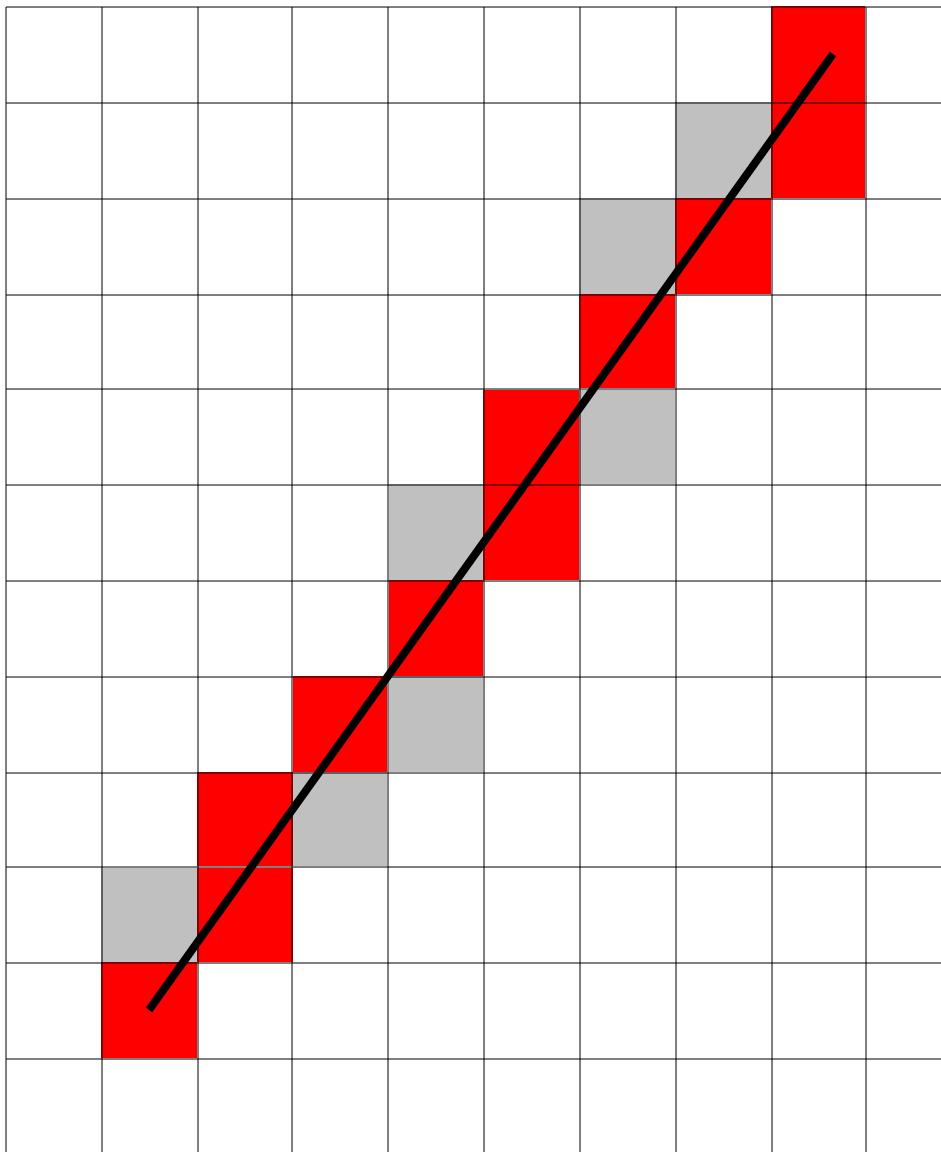
```
   $j += sign_y$ 
   $t_{min} = t_{next\_h}$ 
   $t_{next\_h} += dt_h$ 
```

(dir_x, dir_y)

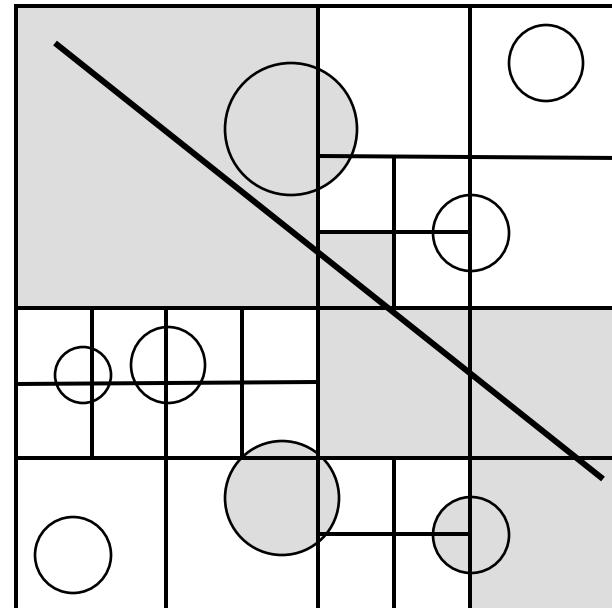
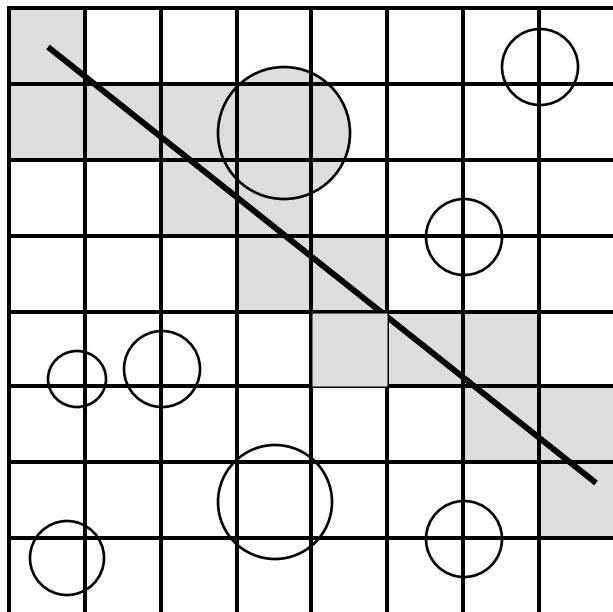


if ($dir_x > 0$) $sign_x = 1$ else $sign_x = -1$
if ($dir_y > 0$) $sign_y = 1$ else $sign_y = -1$

What's the next cell?

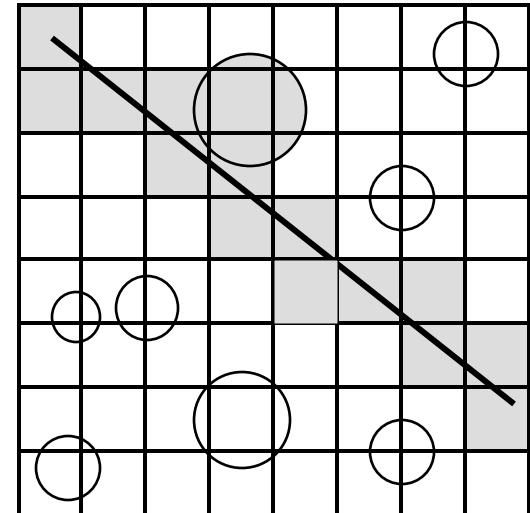


Uniform vs. Adaptive Subdivision



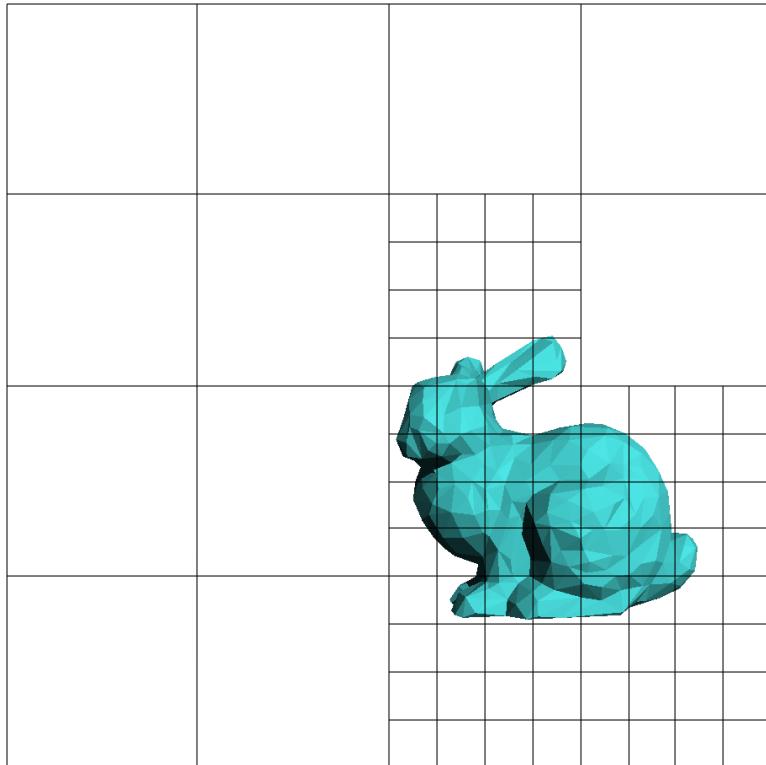
Regular Grid Discussion

- Advantages?
 - easy to construct
 - easy to traverse
- Disadvantages?
 - may be only sparsely filled
 - geometry may still be clumped

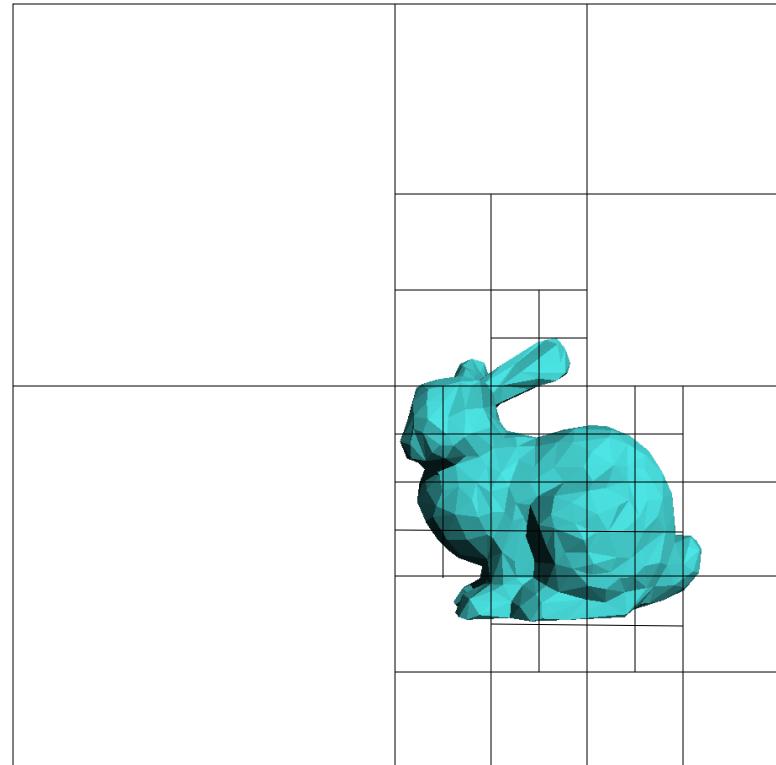


Adaptive Grids

- Subdivide until each cell contains no more than n elements, or maximum depth d is reached



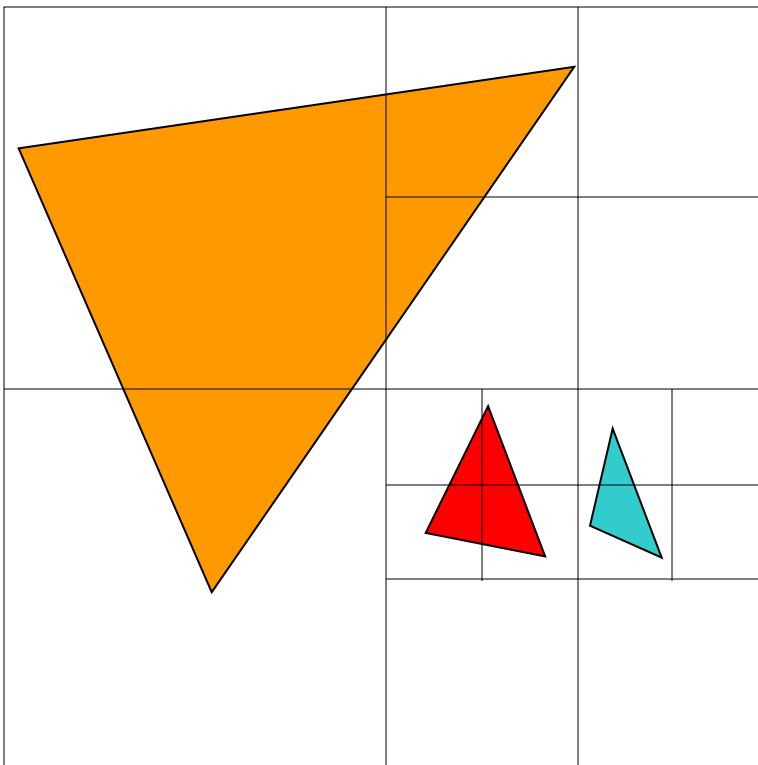
Nested Grids



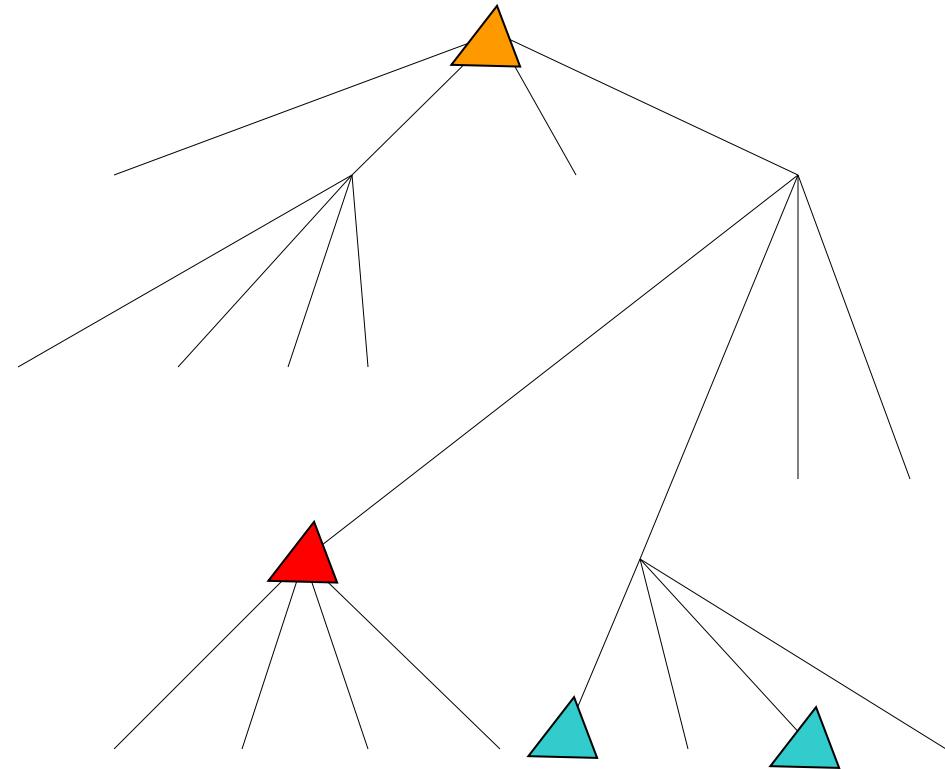
Octree/(Quadtree)

Primitives in an Adaptive Grid

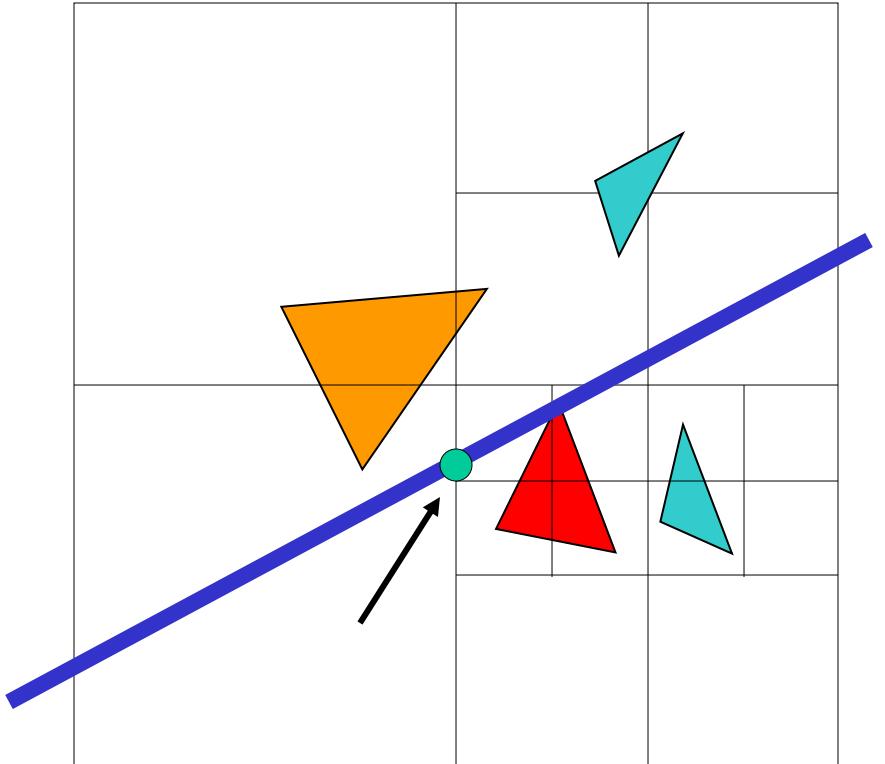
- Can live at intermediate levels, or
be pushed to lowest level of grid



Octree/(Quadtree)

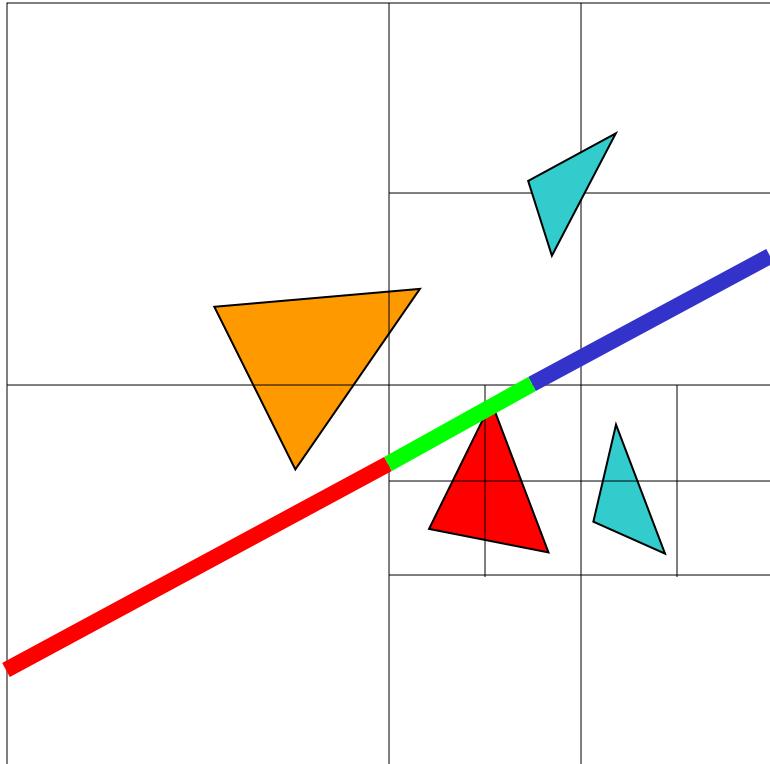


Bottom Up traversal



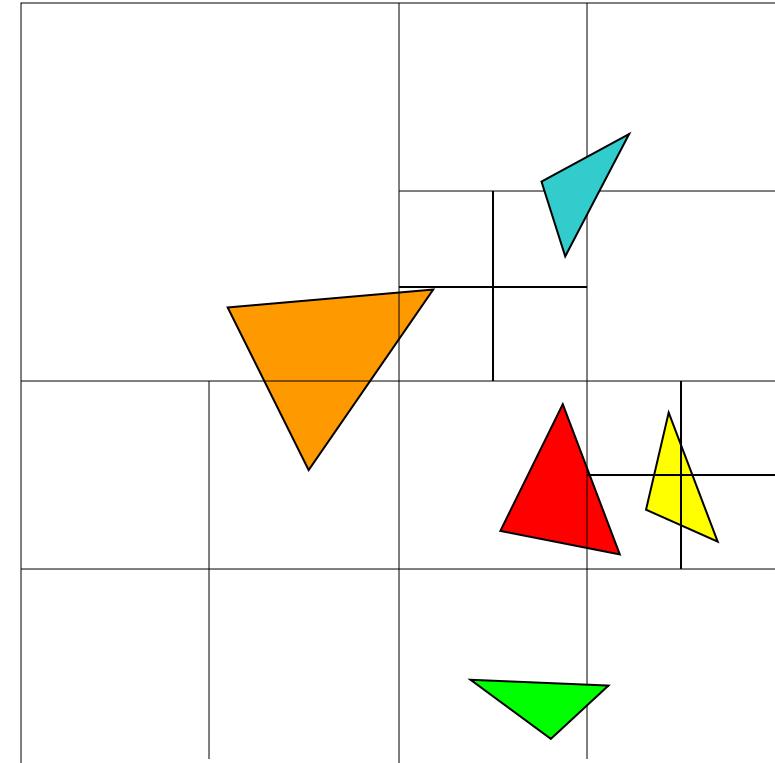
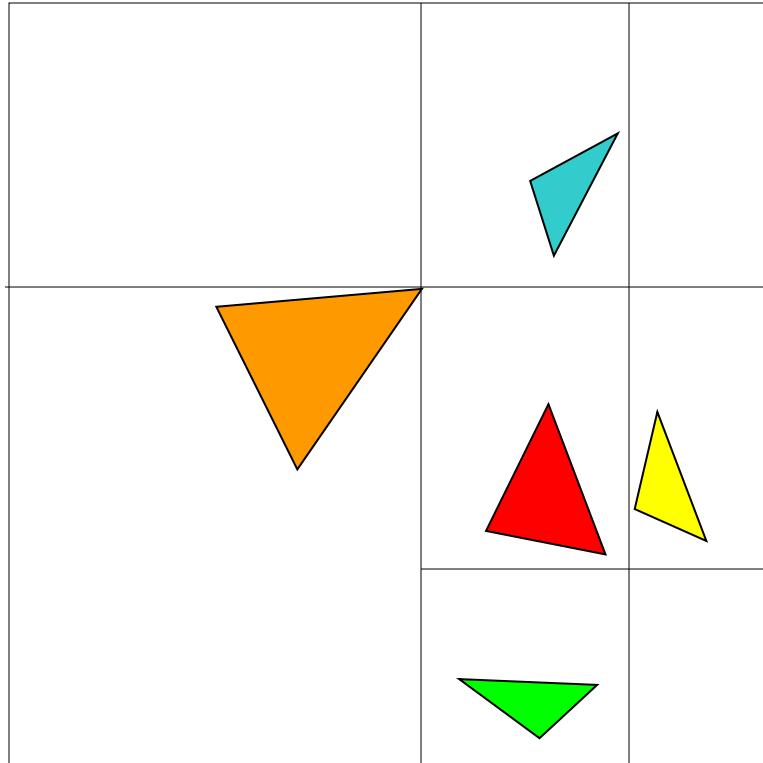
Step from cell to cell.
Intersect current cell and add an epsilon into the next cell.
Then search for the cell in the tree.
A naïve search starts from the root.
Otherwise, try an intelligent guess...

Top down traversal

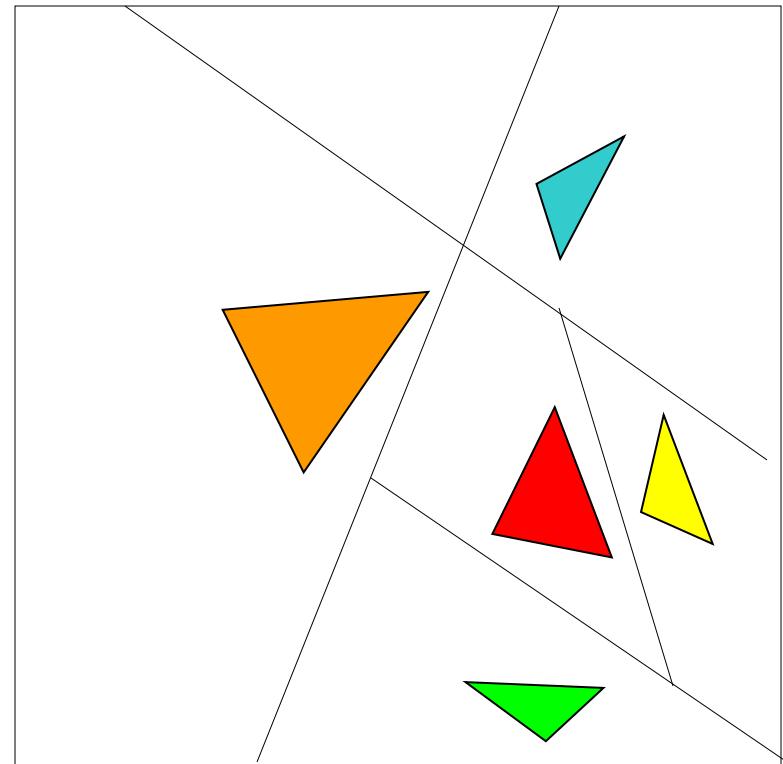
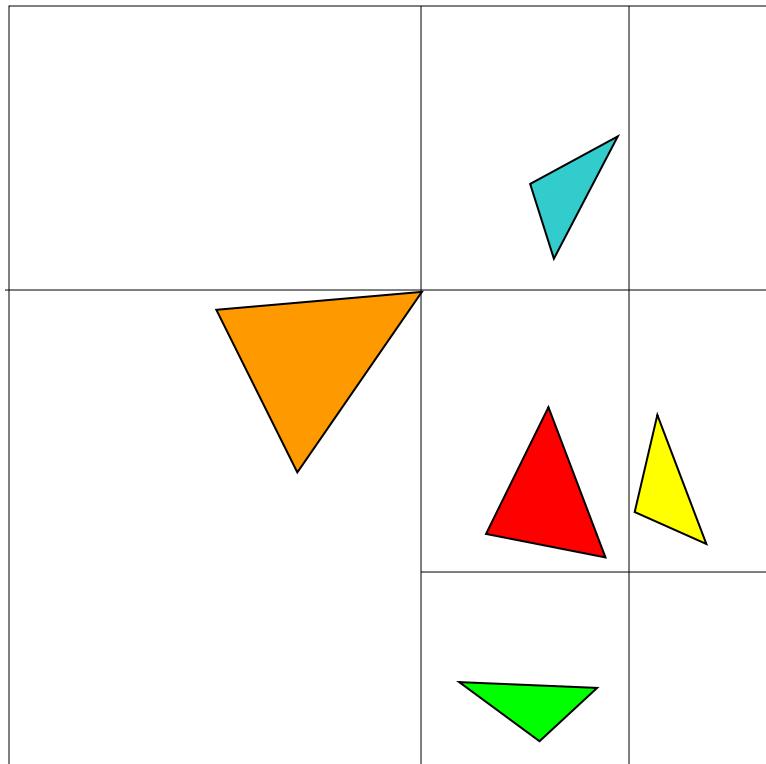


Split ray into sub-segments and traverse each segment recursively.

Kd-trees vs. Quad-tree



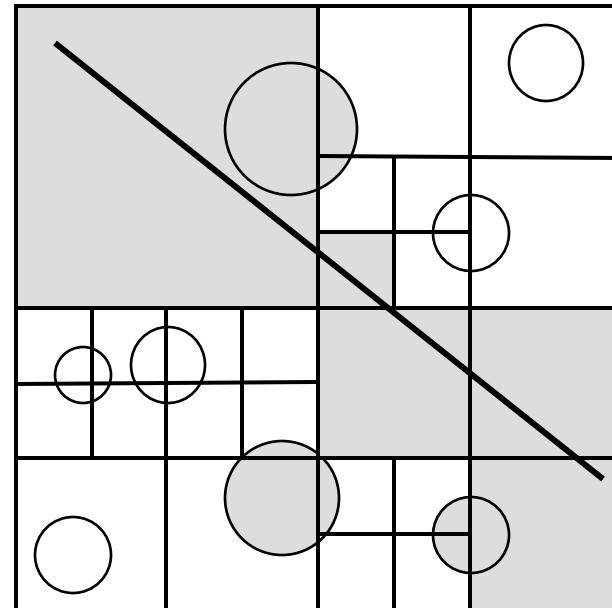
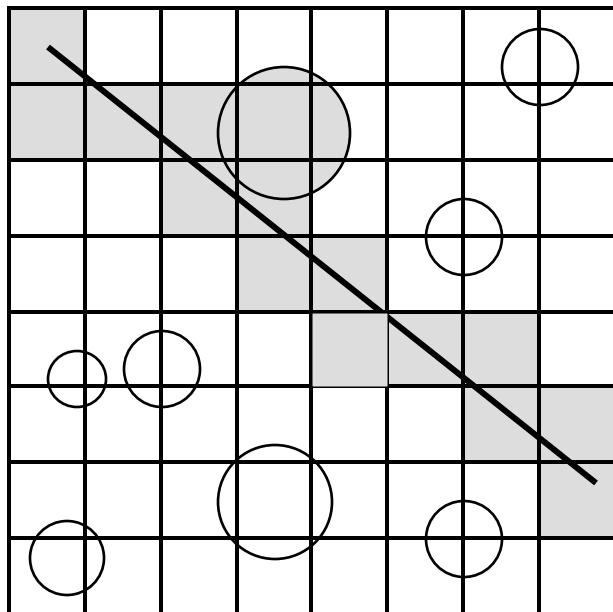
Kd-trees vs. BSP-tree



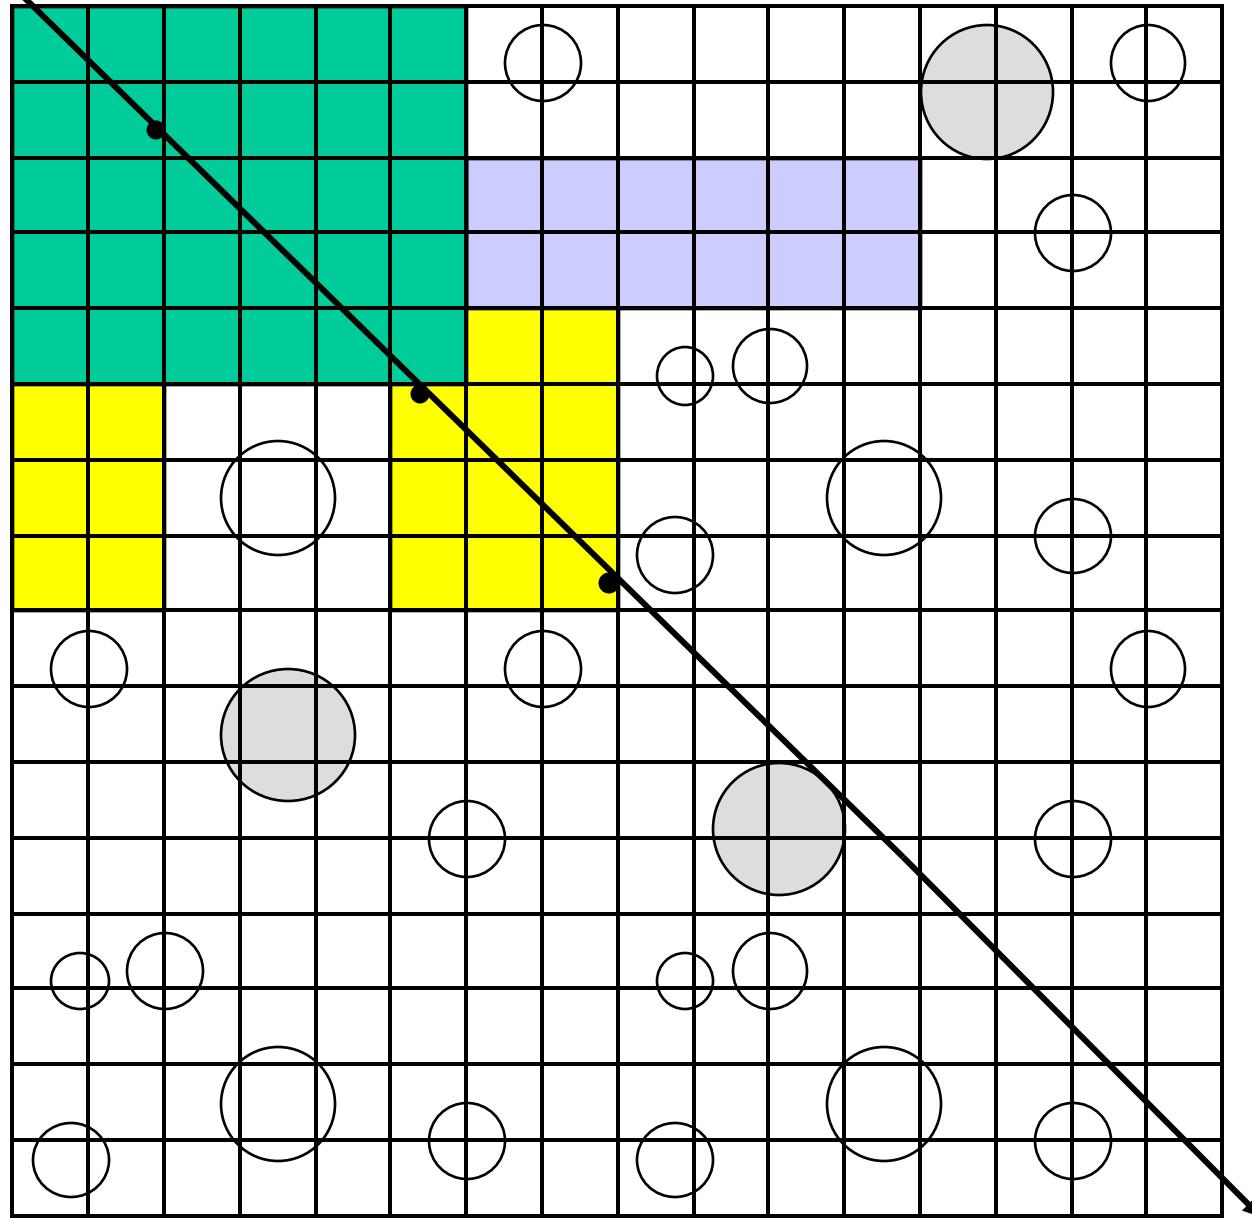
Adaptive Spatial Subdivision

- Disadvantages of uniform subdivision:
 - requires a lot of space
 - traversal of empty regions of space can be slow
 - not suitable for “teapot in a stadium” scenes
- Solution: use a hierarchical adaptive spatial subdivision data structure
 - octrees
 - BSP-trees
- Given a ray, perform a depth-first traversal of the tree. Again, can stop once an intersection has been found.

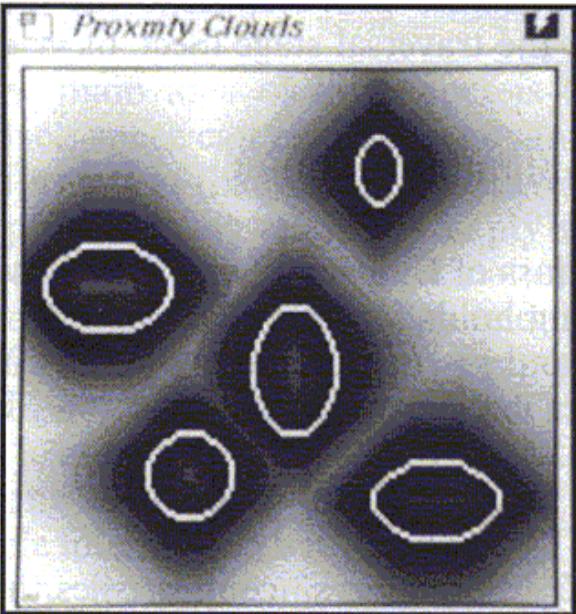
Uniform vs. Adaptive Subdivision



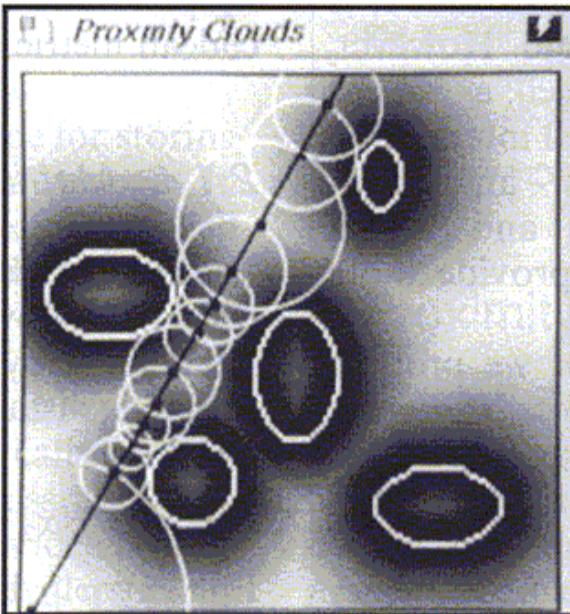
Macro-regions



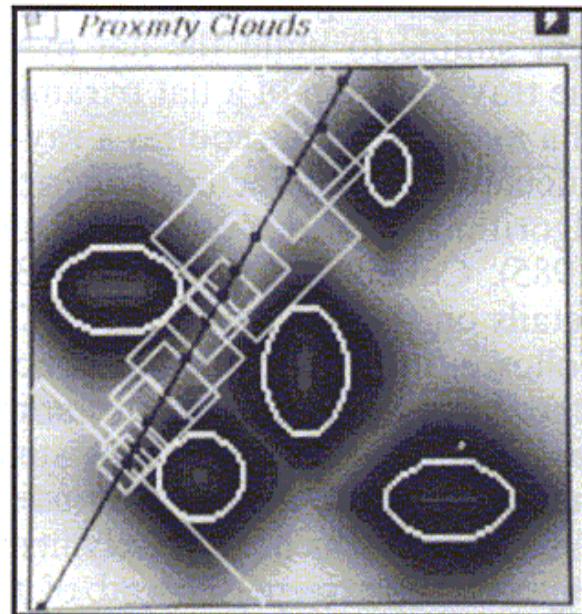
Proximity Clouds



3



4 a



b

Parallel/Distributed RT

- Two main approaches:
 - Each processor is in charge of tracing a subset of the rays. Requires a shared memory architecture, replication of the scene database, or transmission of objects between processors on demand.
 - Each processor is in charge of a subset of the scene (either in terms of space, or in terms of objects). Requires processors to transmit rays among themselves.

Directional Techniques

- Light buffer: accelerates shadow rays.
 - Discretize the space of directions around each light source using the *direction cube*
 - In each cell of the cube store a sorted list of objects visible from the light source through that cell
 - Given a shadow ray locate the appropriate cell of the direction cube and test the ray with the objects on its list

Directional Techniques

- Ray classification (Arvo and Kirk 87):
 - Rays in 3D have 5 degrees of freedom: (x, y, z, θ, ϕ)
 - Rays coherence: rays belonging to the same small 5D neighborhood are likely to intersect the same set of objects.
 - Partition the 5D space of rays into a collection of 5D hypercubes, each containing a list of objects.
 - Given a ray, find the smallest containing 5D hypercube, and test the ray against the objects on the list.
 - For efficiency, the hypercubes are arranged in a hierarchy: a 5D analog of the 3D octree. This data structure is constructed in a lazy fashion.